

COMPUTER VISION LECTURE 7 – LINE DETECTION

Prof. Dr. Francesco Maurelli
2018-09-25

WHAT WE WILL LEARN TODAY

1. Recap Edge detection
2. Hough Transform
3. RANSAC
4. Intro to Features



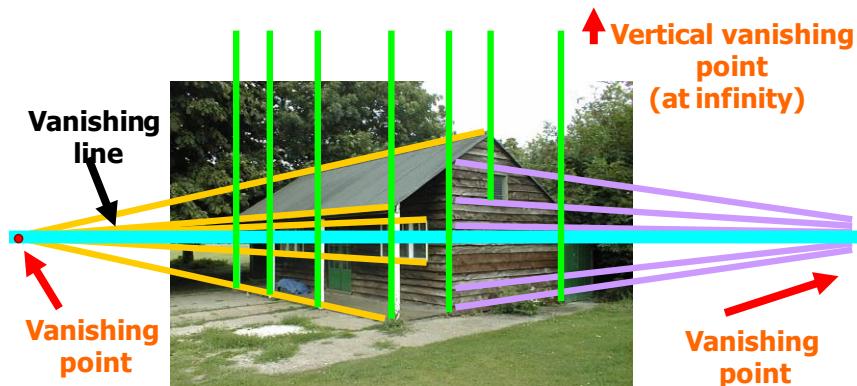
- **Goal:** Identify sudden changes (discontinuities) in an image
 - Intuitively, most semantic and shape information from the image can be encoded in the edges
 - More compact than pixels
- **Ideal:** artist's line drawing (but artist is also using object-level knowledge)



WHY DO WE CARE ABOUT EDGES?

1. Extract information, recognize objects

1. Recover geometry and viewpoint



ORIGINS OF EDGES



surface normal discontinuity

depth discontinuity

surface color discontinuity

illumination discontinuity

Backward

$$\frac{df}{dx} = f(x) - f(x-1) = f'(x)$$

Forward

$$\frac{df}{dx} = f(x) - f(x+1) = f'(x)$$

Central

$$\frac{df}{dx} = f(x+1) - f(x-1) = f'(x)$$

Backward filter: [0 1 -1]

$$f(x) - f(x-1) = f'(x)$$

Forward: [-1 1 0]

$$f(x) - f(x+1) = f'(x)$$

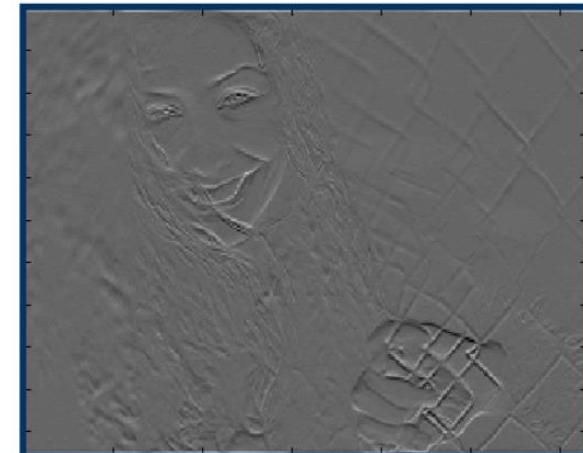
Central: [1 0 -1]

$$f(x+1) - f(x-1) = f'(x)$$

3X3 IMAGE GRADIENT FILTERS

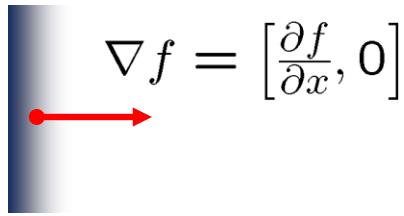
$$\frac{1}{3} \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

$$\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$



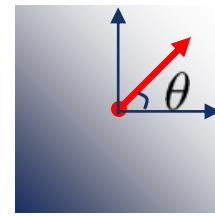
$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

1. The gradient of an image:


$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient vector points in the direction of most rapid increase in intensity

The gradient direction is given by $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$

The *edge strength* is given by the gradient magnitude

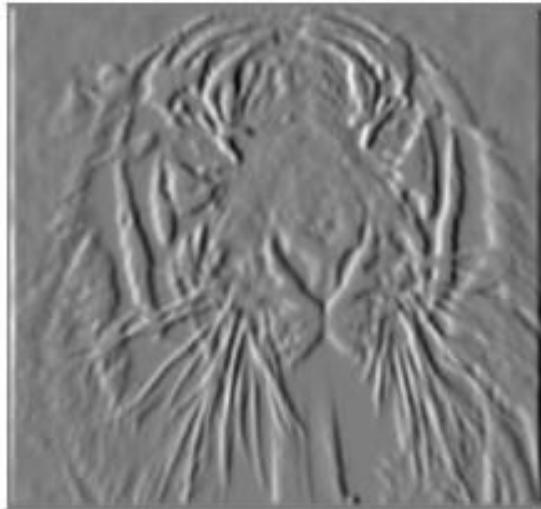
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

FINITE DIFFERENCES: EXAMPLE

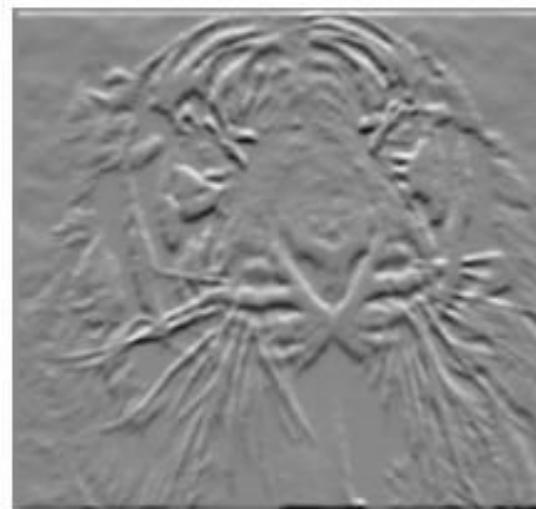
Original
Image



x-direction



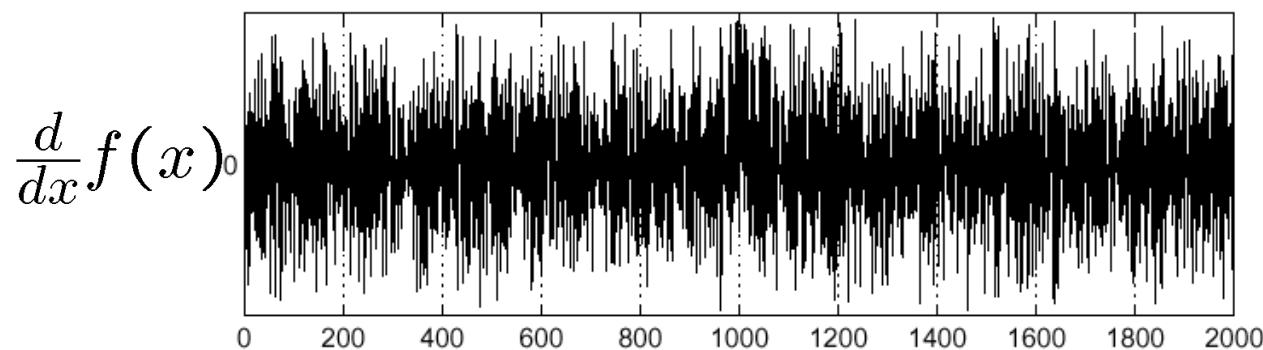
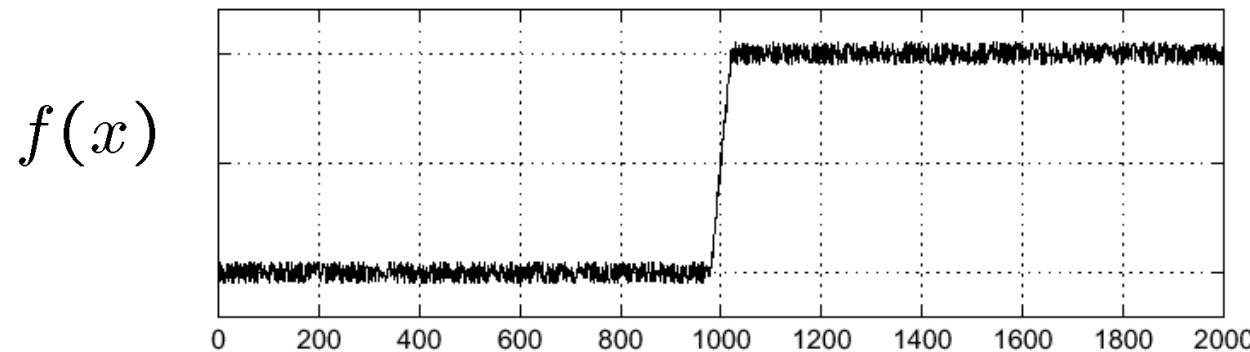
Gradient
magnitude



y-direction

Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal



Where is the edge?

SMOOTHING WITH DIFFERENT FILTERS

Mean



Gaussian



Median



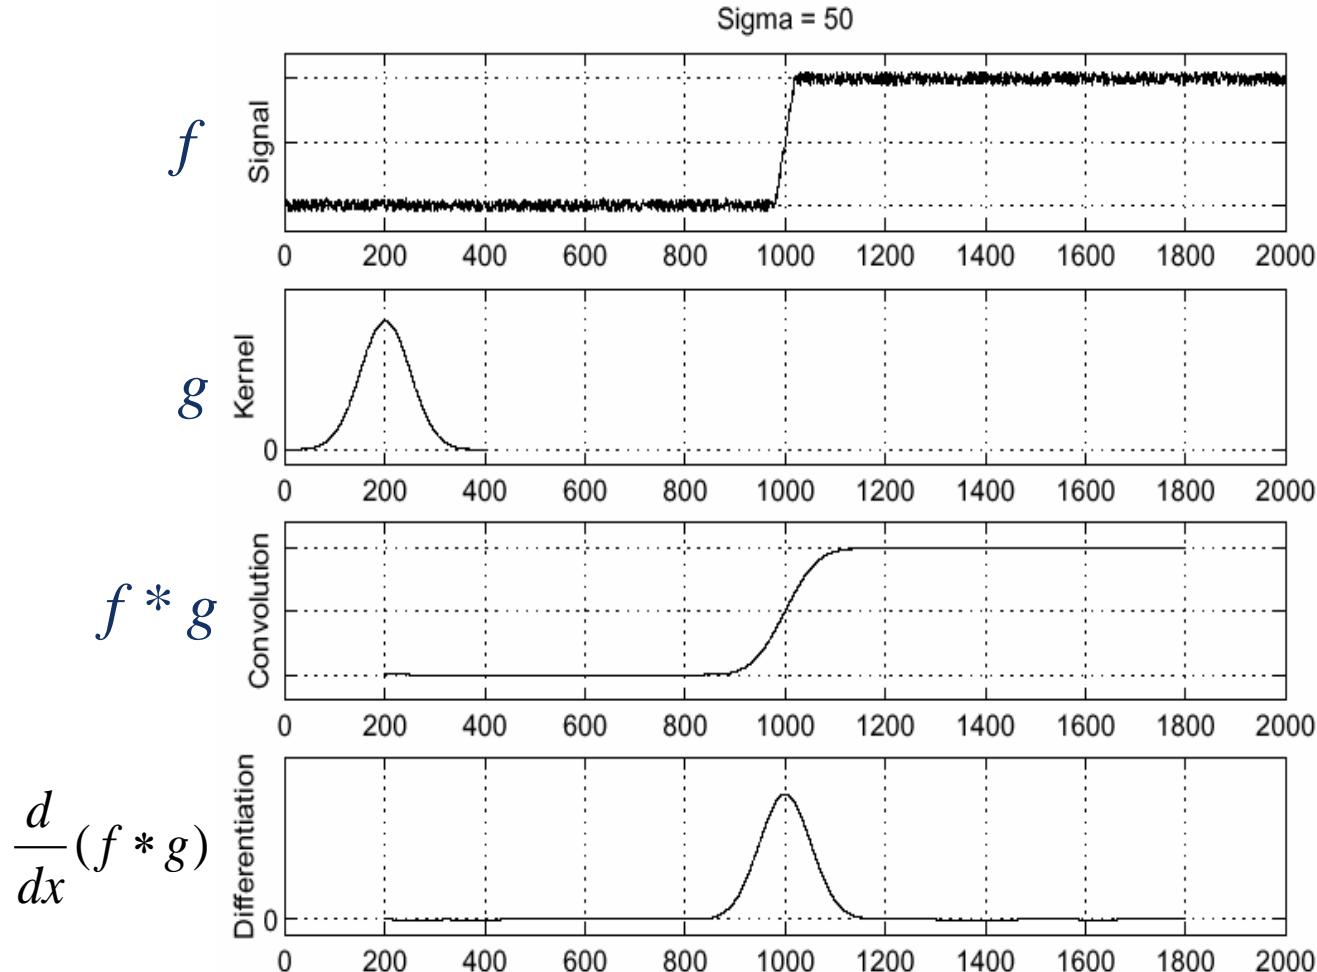
3x3

5x5

7x7



SOLUTION: SMOOTH FIRST

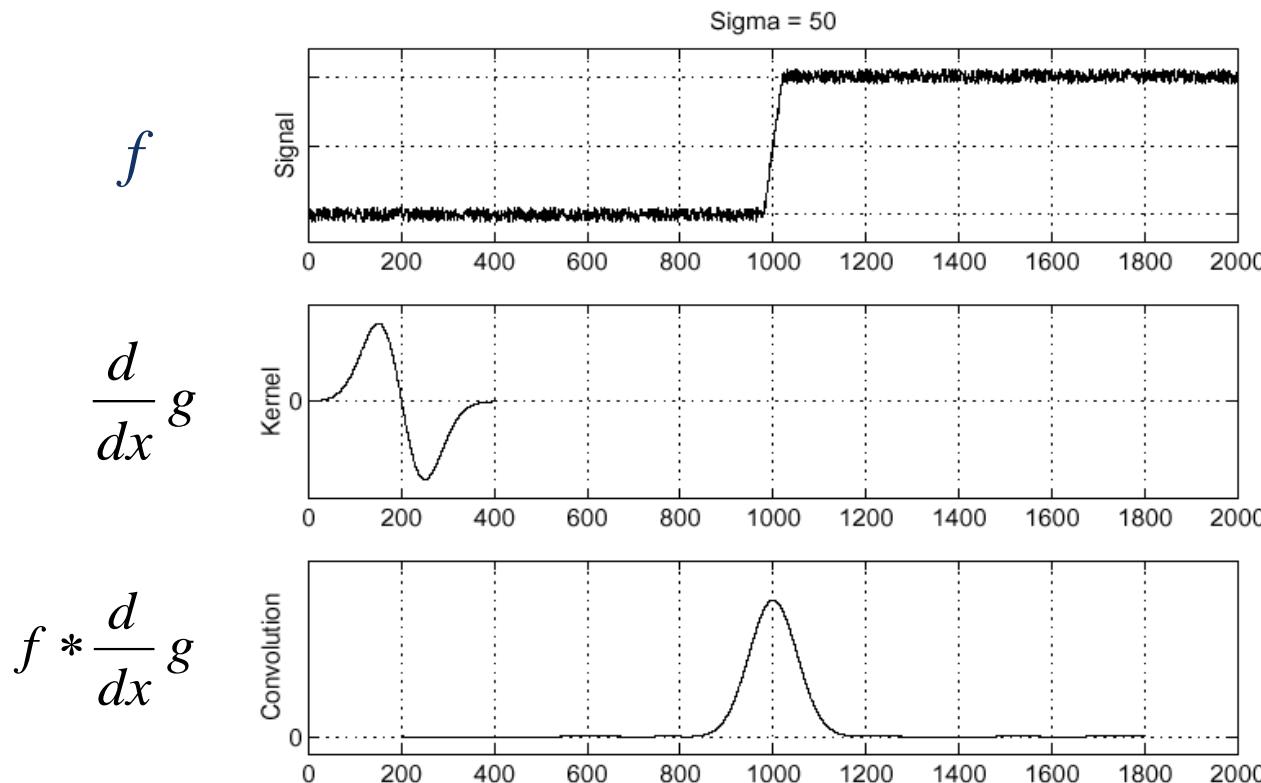


to find edges, look for peaks in $\frac{d}{dx}(f * g)$

DERIVATIVE THEOREM OF CONVOLUTION

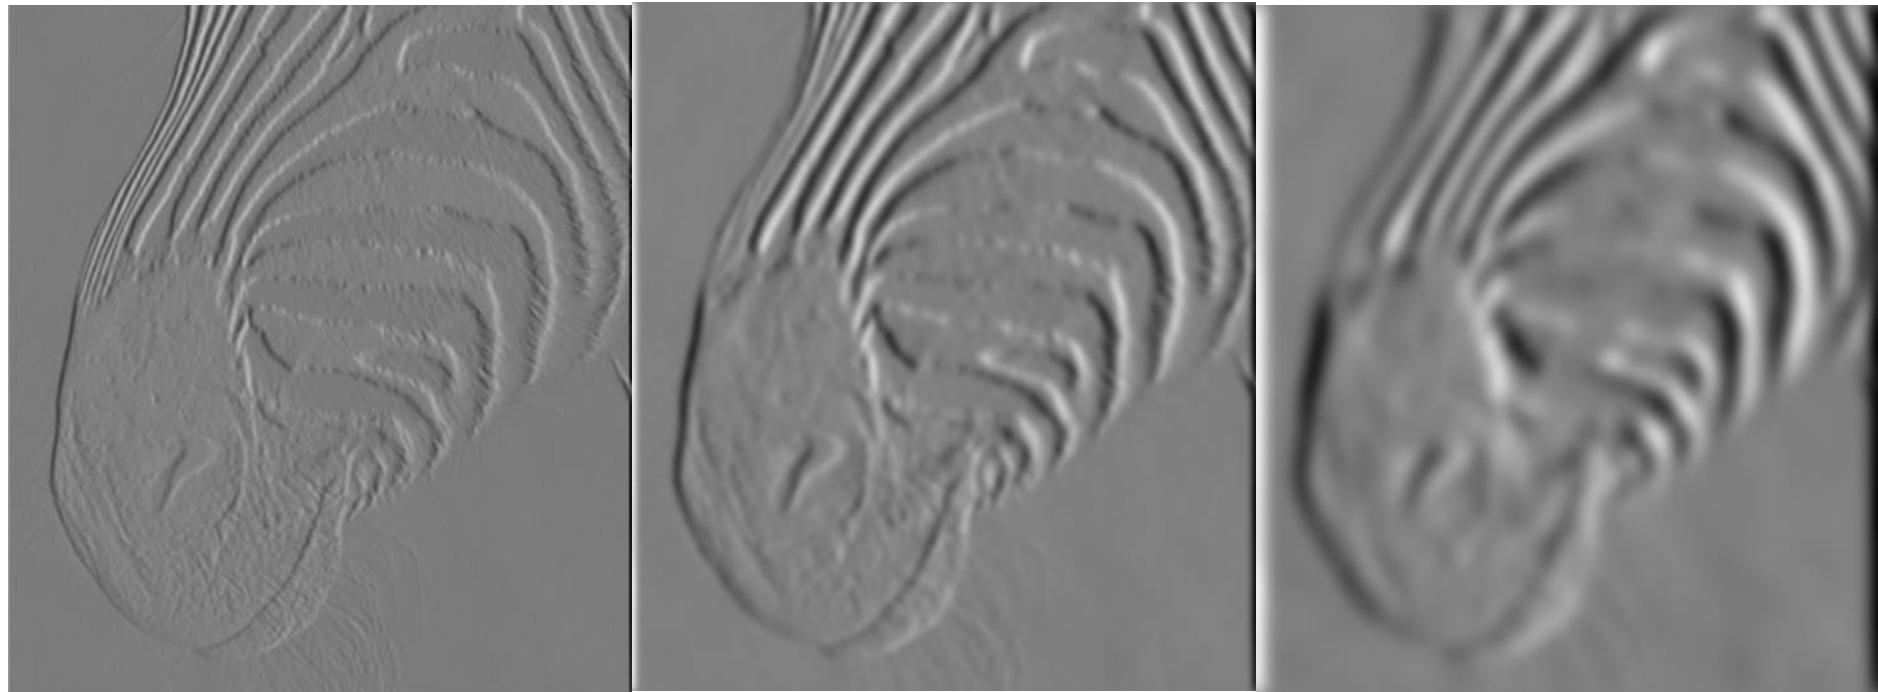
- This theorem gives us a very useful property:
- This saves us one operation:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$



Source: S. Seitz

TRADEOFF BETWEEN SMOOTHING AT DIFFERENT SCALES



1 pixel

3 pixels

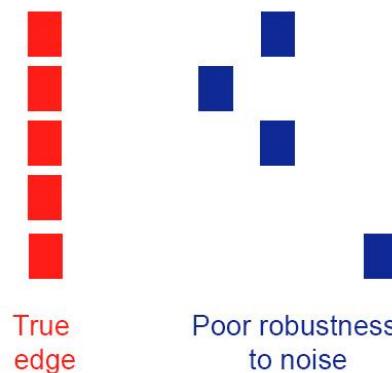
7 pixels

Smoothed derivative removes noise, but blurs edge.
Also finds edges at different “scales”.

DESIGNING AN EDGE DETECTOR

Criteria for an “optimal” edge detector:

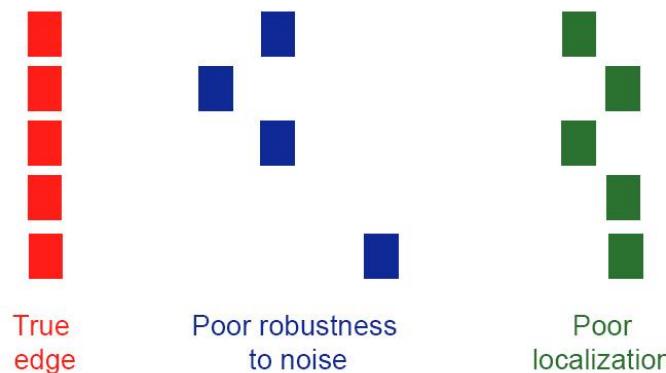
- **Good detection:** the optimal detector must minimize the probability of false positives (detecting spurious edges caused by noise), as well as that of false negatives (missing real edges)



DESIGNING AN EDGE DETECTOR

Criteria for an “optimal” edge detector:

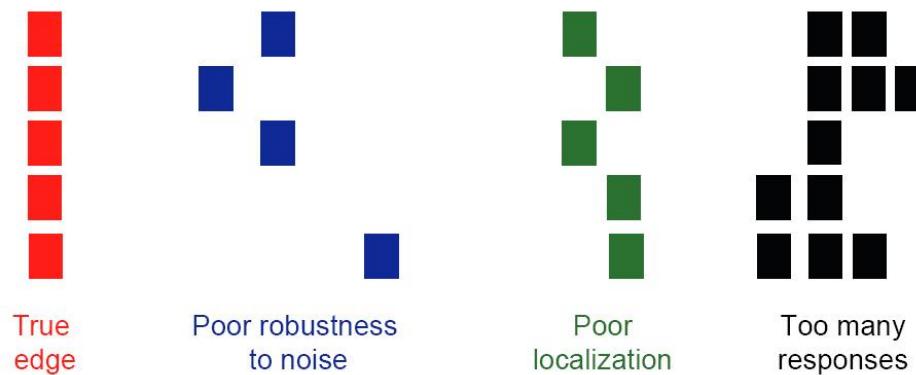
- **Good detection:** the optimal detector must minimize the probability of false positives (detecting spurious edges caused by noise), as well as that of false negatives (missing real edges)
- **Good localization:** the edges detected must be as close as possible to the true edges



DESIGNING AN EDGE DETECTOR

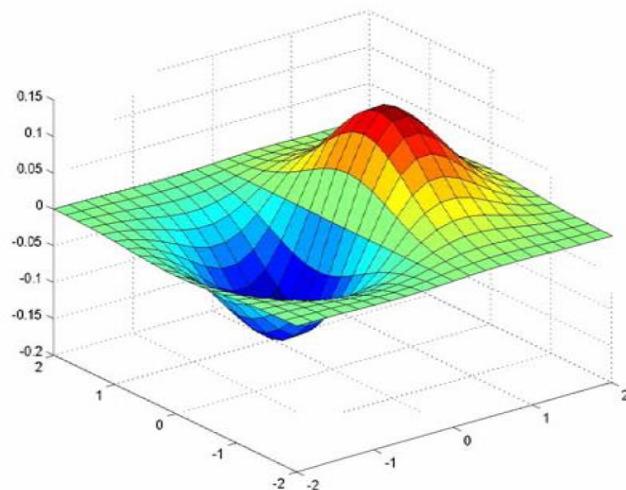
Criteria for an “optimal” edge detector:

- **Good detection:** the optimal detector must minimize the probability of false positives (detecting spurious edges caused by noise), as well as that of false negatives (missing real edges)
- **Good localization:** the edges detected must be as close as possible to the true edges
- **Single response:** the detector must return one point only for each true edge point; that is, minimize the number of local maxima around the true edge

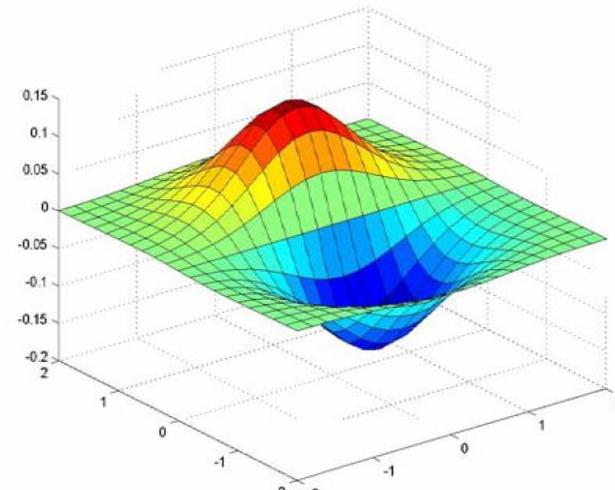


1. Suppress Noise
2. Compute gradient magnitude and direction
3. Apply Non-Maximum Suppression
 - Assures minimal response
4. Use hysteresis and connectivity analysis to detect edges

DERIVATIVE OF GAUSSIAN FILTER

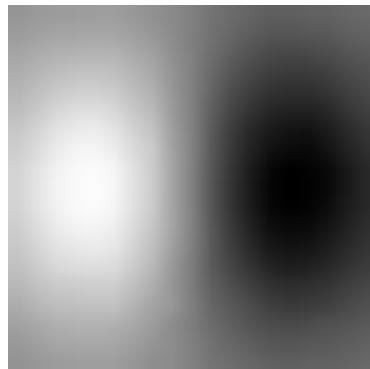


x-direction

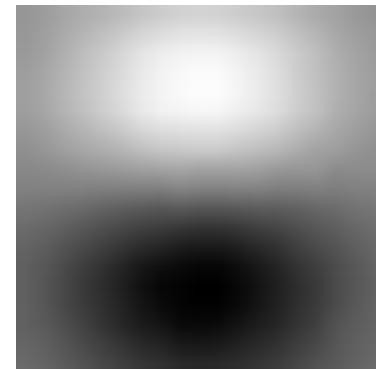


y-direction

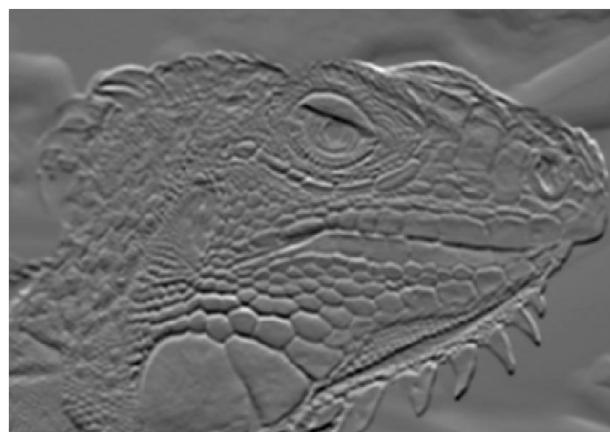
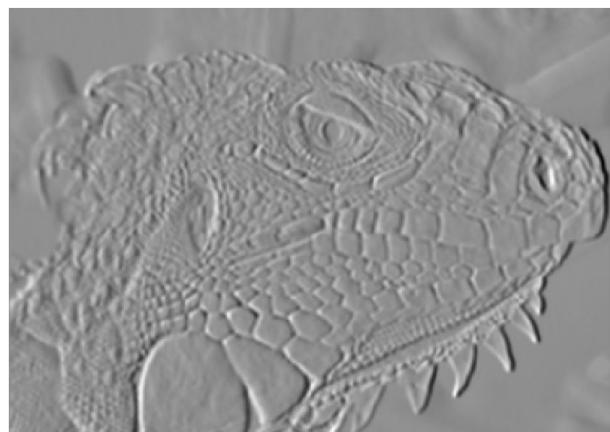
$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$



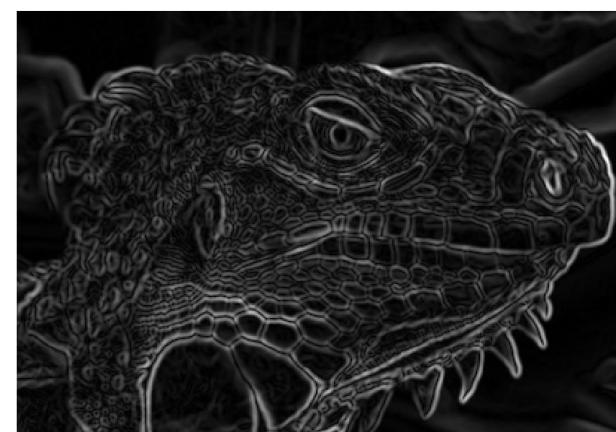
$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$



COMPUTE GRADIENTS



X-Derivative of Gaussian



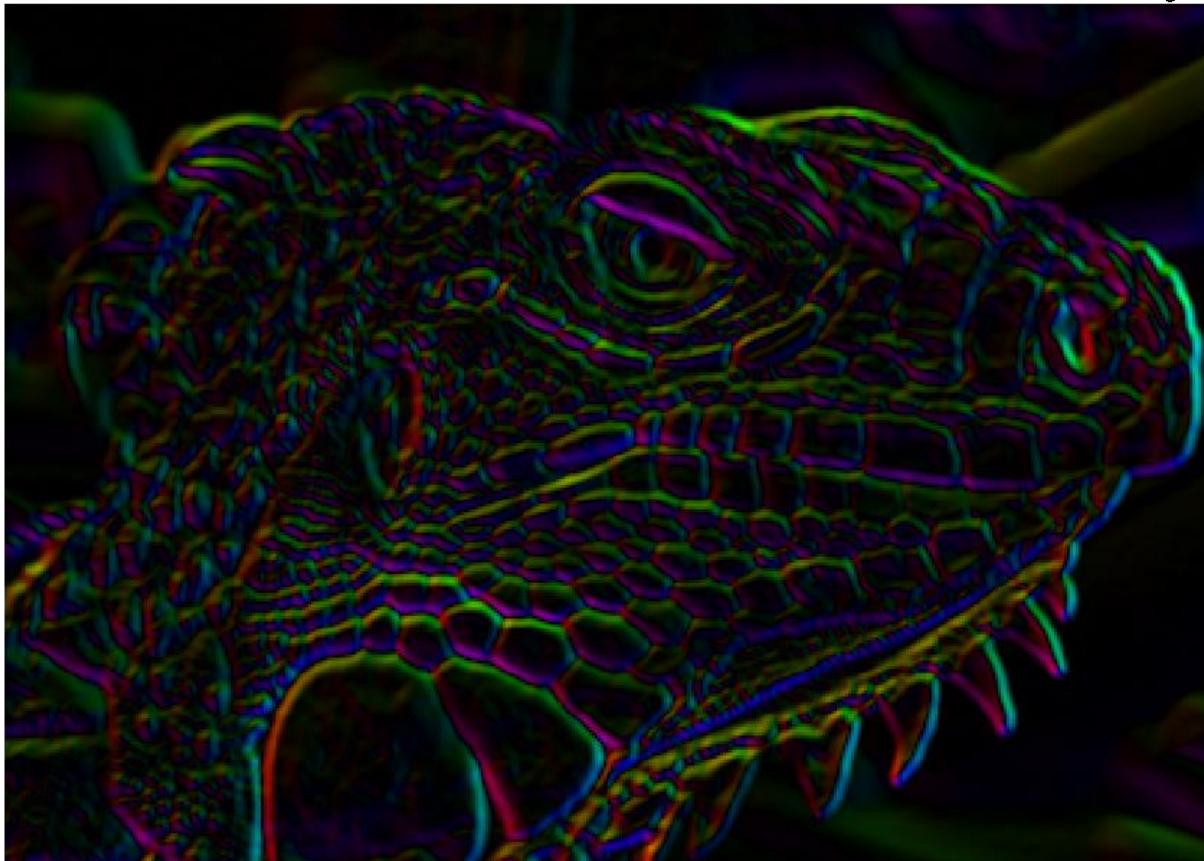
Y-Derivative of Gaussian

Gradient Magnitude

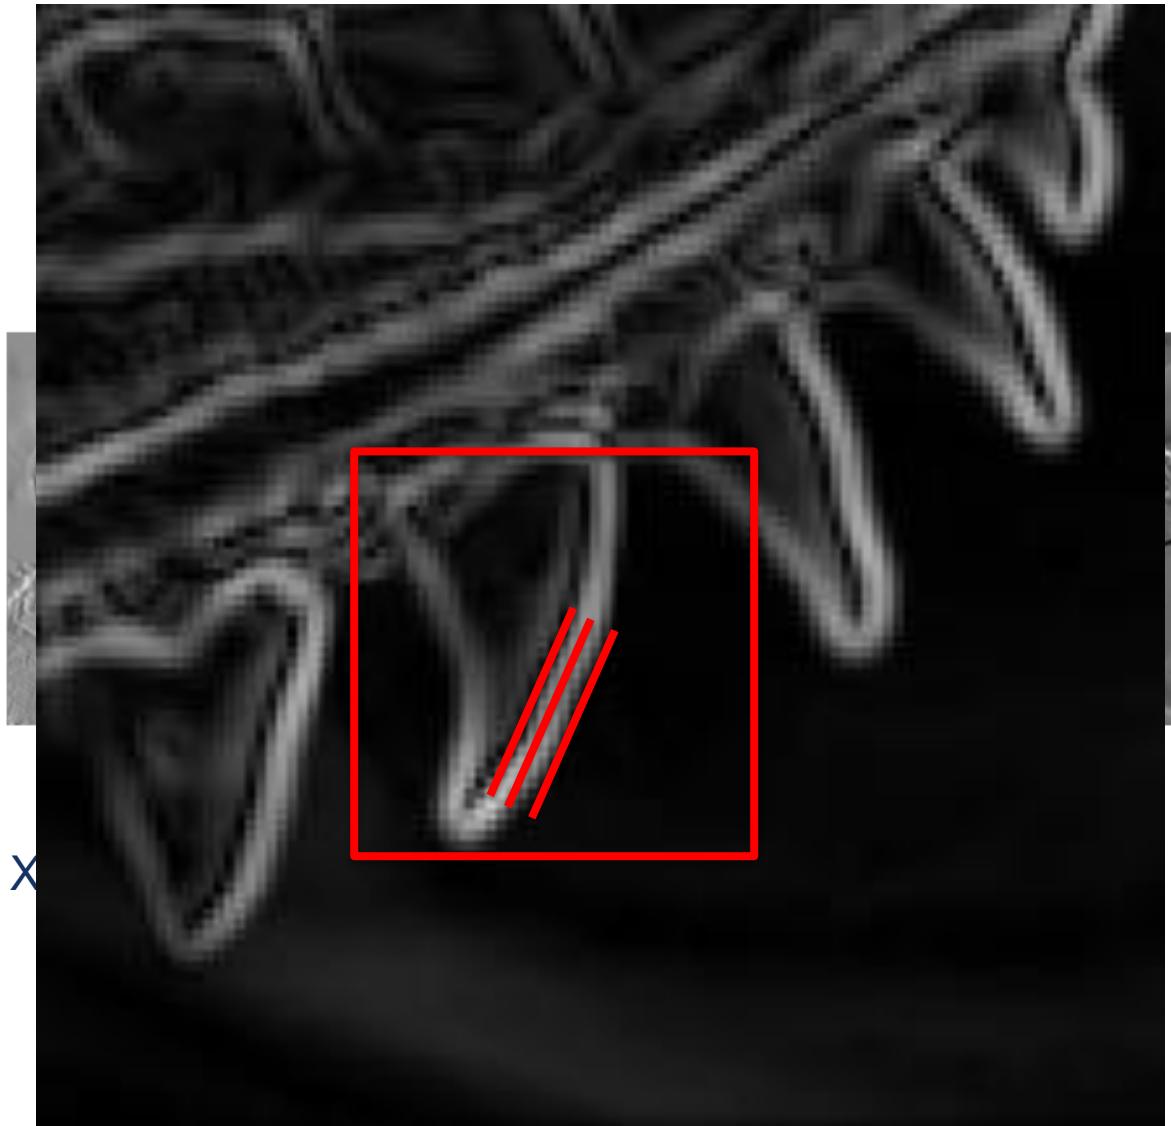
Source: J. Hayes

GET ORIENTATION AT EACH PIXEL

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

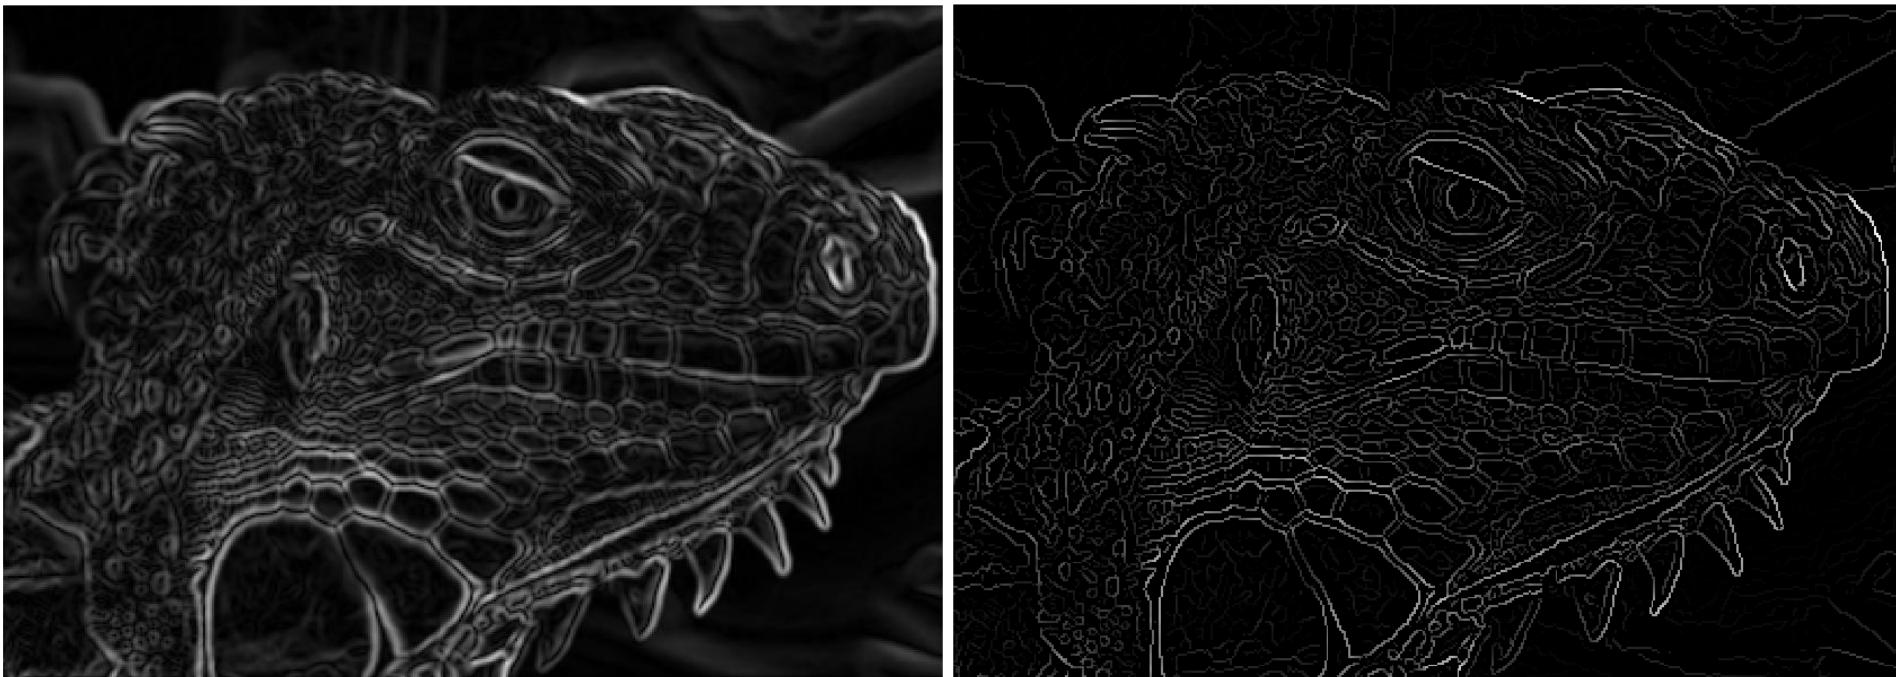


COMPUTE GRADIENTS



Gradient Magnitude

NON-MAX SUPPRESSION



Before

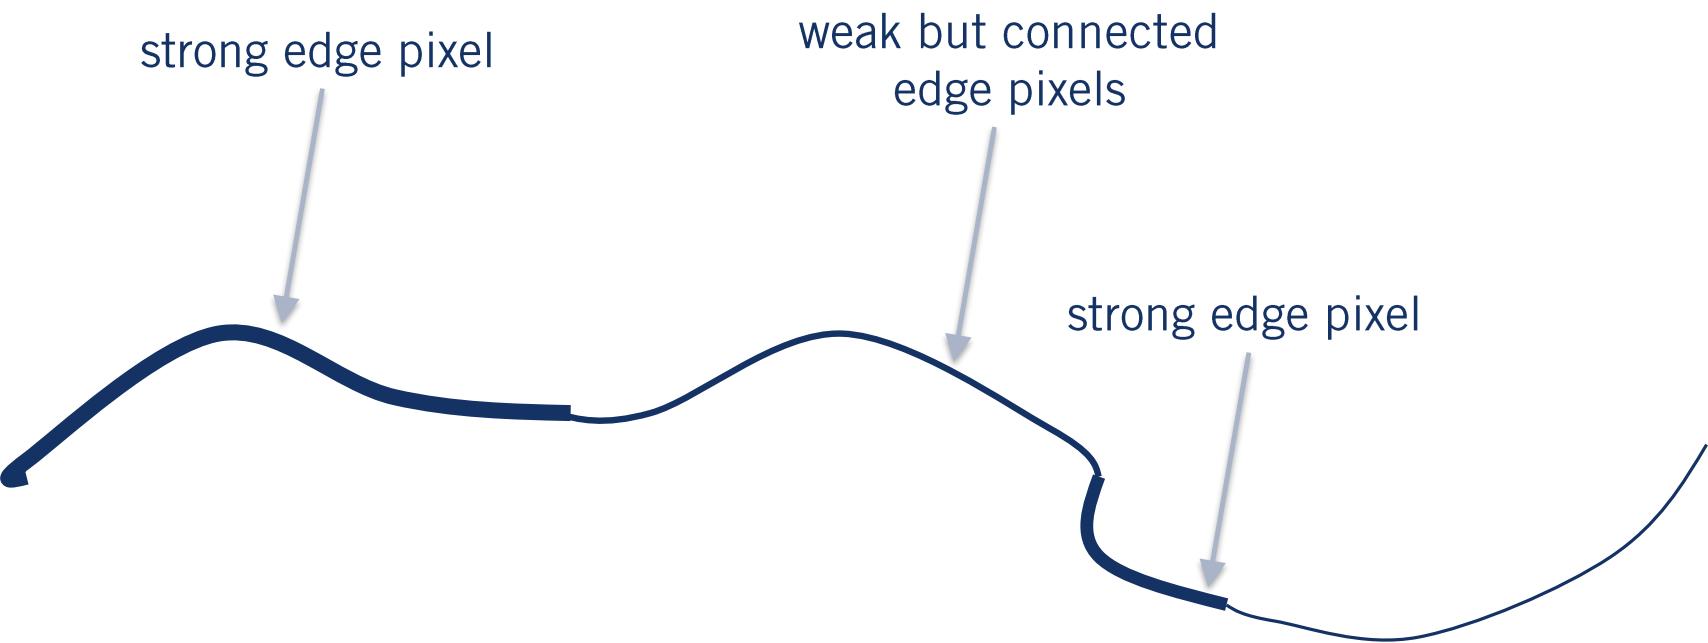
After

If the gradient at a pixel is

- 1.above High, declare it as an ‘strong edge pixel’
- 2.below Low, declare it as a “non-edge-pixel”
- 3.between Low and High

- Consider its neighbors iteratively then declare it an “edge pixel” if it is connected to a ‘strong edge pixel’ directly or via pixels between Low and High

HYSTERESIS THRESHOLDING



Source: S. Seitz

FINAL CANNY EDGES





Original Image

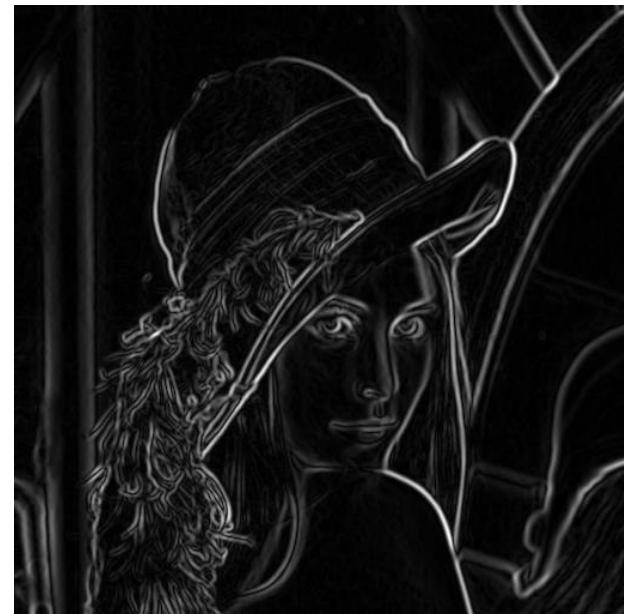
CANNY EDGE DETECTOR



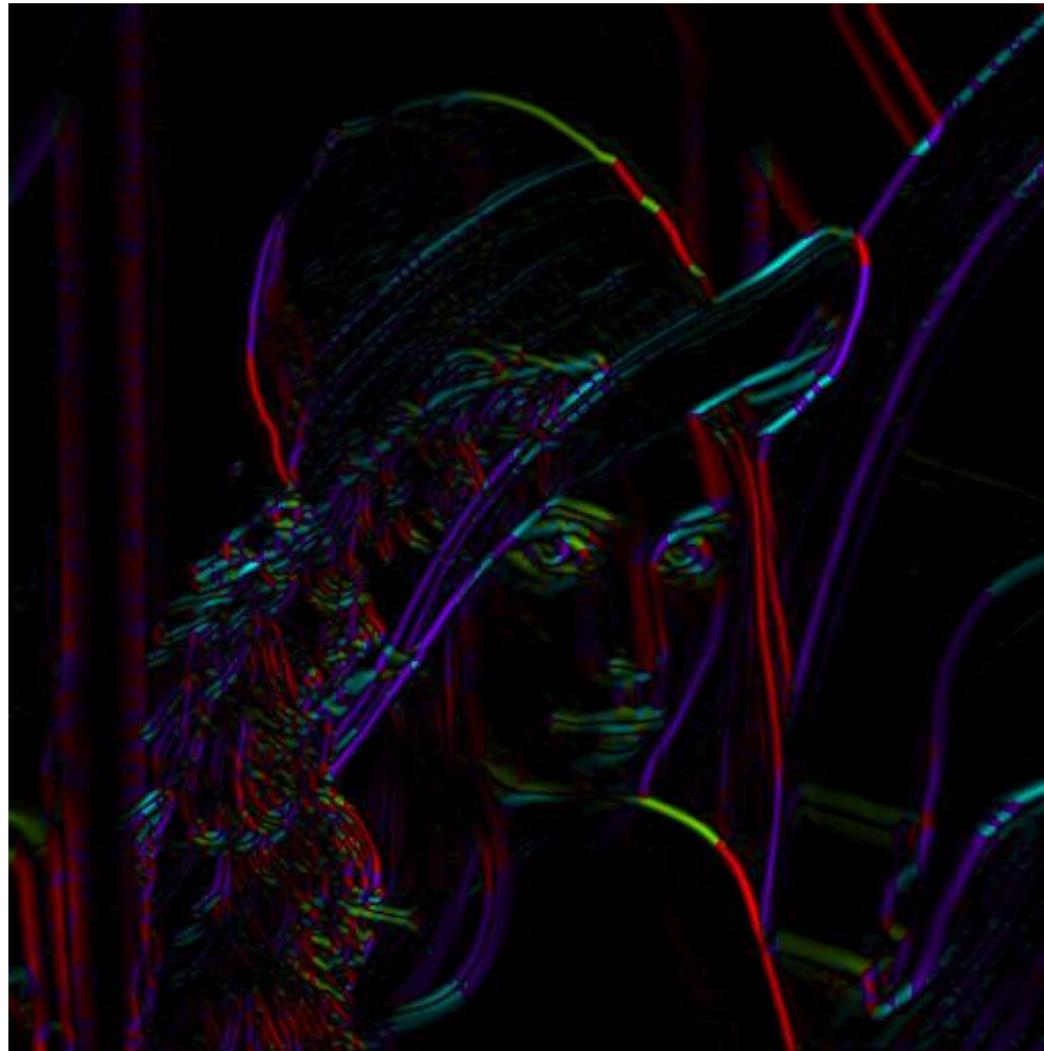
X-Derivative



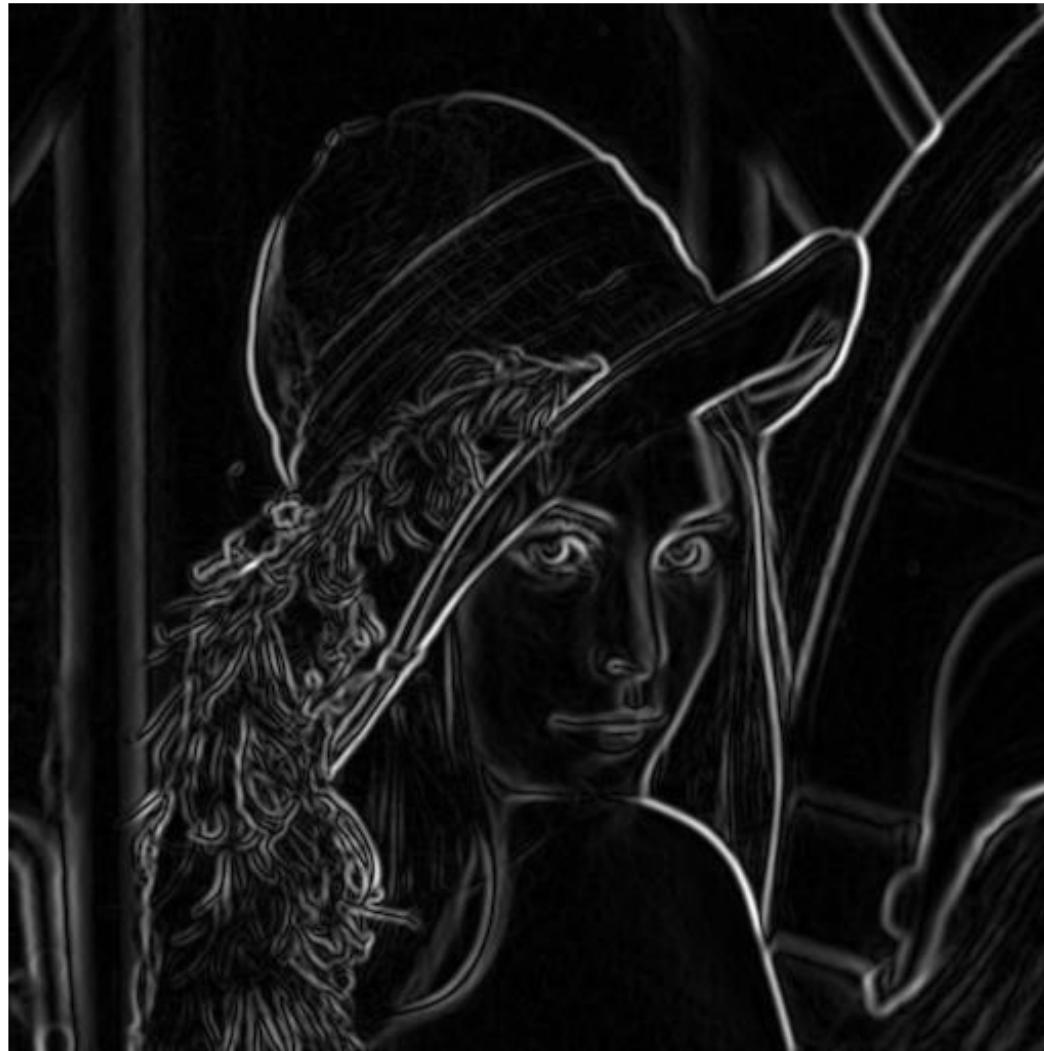
Y-Derivative



Gradient Magnitude



Orientation at each pixel



Before non-max suppression



After non-max suppression



Hysteresis thresholding



Final Canny Edges

EFFECT OF σ (GAUSSIAN KERNEL SPREAD/SIZE)



original



Canny with $\sigma = 1$



Canny with $\sigma = 2$

The choice of σ depends on desired behavior

- large σ detects large scale edges
- small σ detects fine features

- Choose a parametric model to represent a set of features
- Membership criterion is not local
 - Can't tell whether a point belongs to a given model just by looking at that point.
- Three main questions:
 - What model represents this set of features best?
 - Which of several model instances gets which feature?
 - How many model instances are there?
- Computational complexity is important
 - It is infeasible to examine every possible set of parameters and every possible combination of features

LINE FITTING

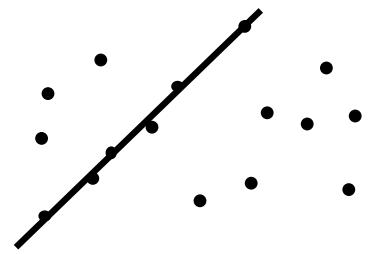
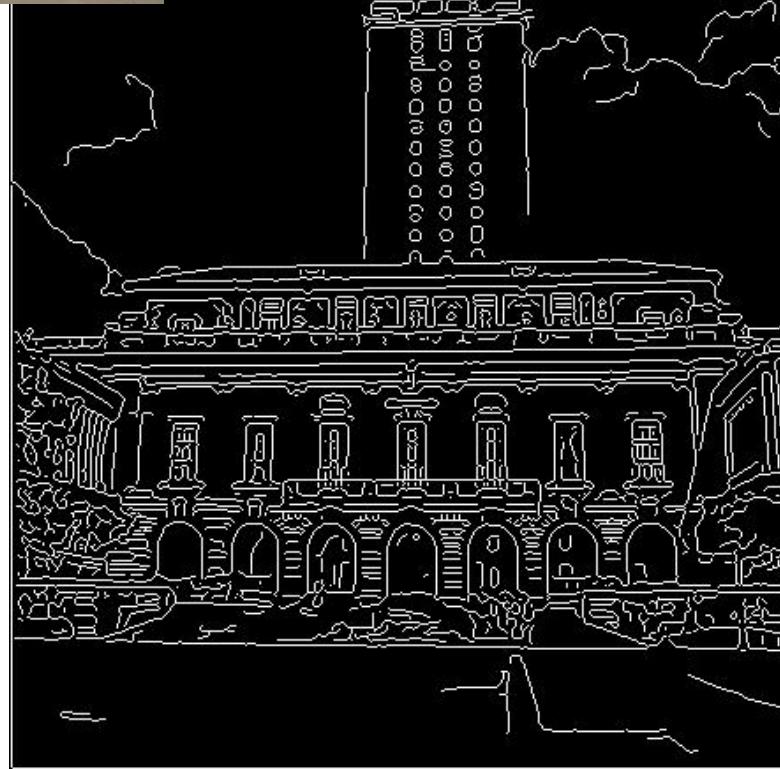
Why?

Many objects characterized by presence of straight lines



Wait, why aren't we done just by running edge detection?

DIFFICULTY OF LINE FITTING

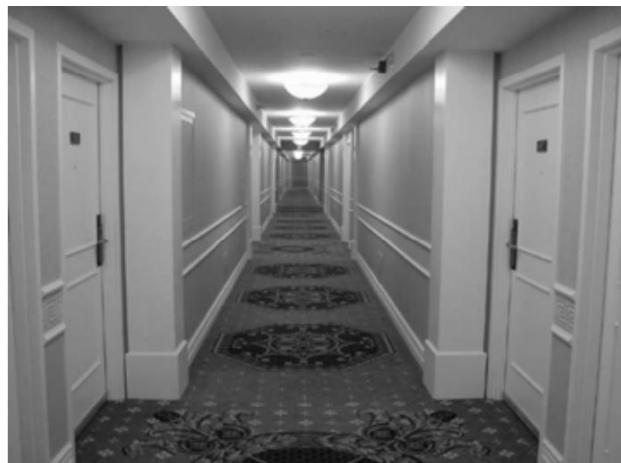


- Extra edge points (clutter), multiple models:
 - Which points go with which line, if any?
- Only some parts of each line detected, and some parts are missing:
 - How to find a line that bridges missing evidence?
- Noise in measured edge points, orientations:
 - How to detect true underlying parameters?

- The Hough transform (HT) can be used to detect lines.
- It was introduced in 1962 (Hough 1962) and first used to find lines in images a decade later (Duda 1972).
- The goal is to find the location of lines in images.
- Caveat: Hough transform can detect lines, circles and other structures ONLY if their parametric equation is known.
- It can give robust detection under noise and partial occlusion

Assume that we have performed some edge detection, and a thresholding of the edge magnitude image.

Thus, we have some pixels that may partially describe the boundary of some objects.



We wish to find sets of pixels that make up straight lines.

Consider a point of known coordinates $(x_j; y_j)$

- There are many lines passing through the point (x_j, y_j) .

Straight lines that pass that point have the form $y_j = a^*x_j + b$

- Common to them is that they satisfy the equation for some set of parameters (a, b)

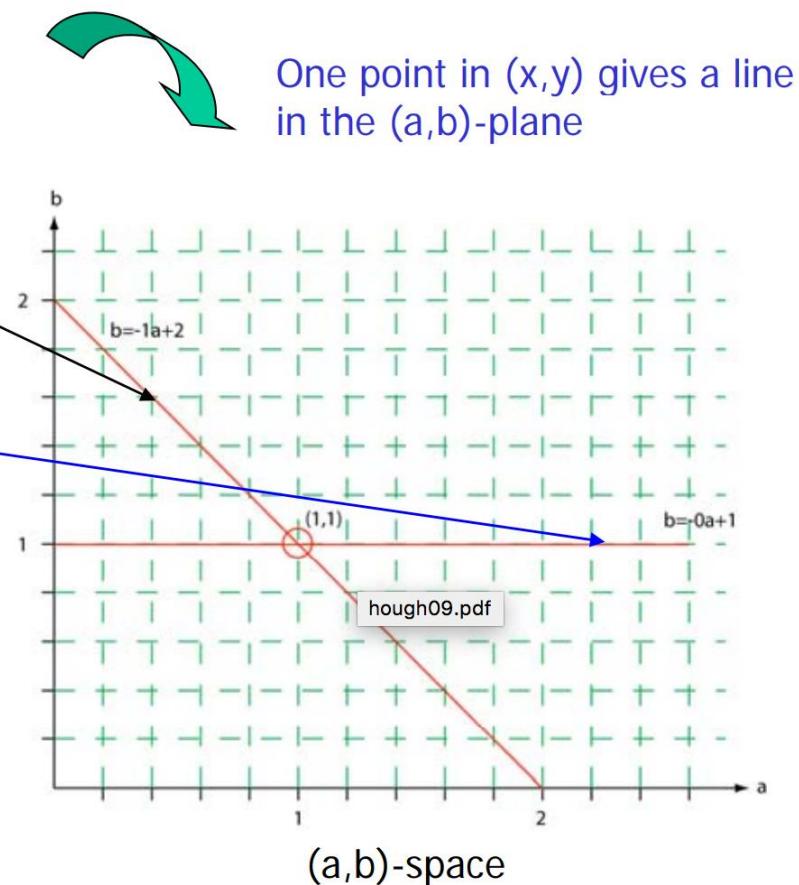
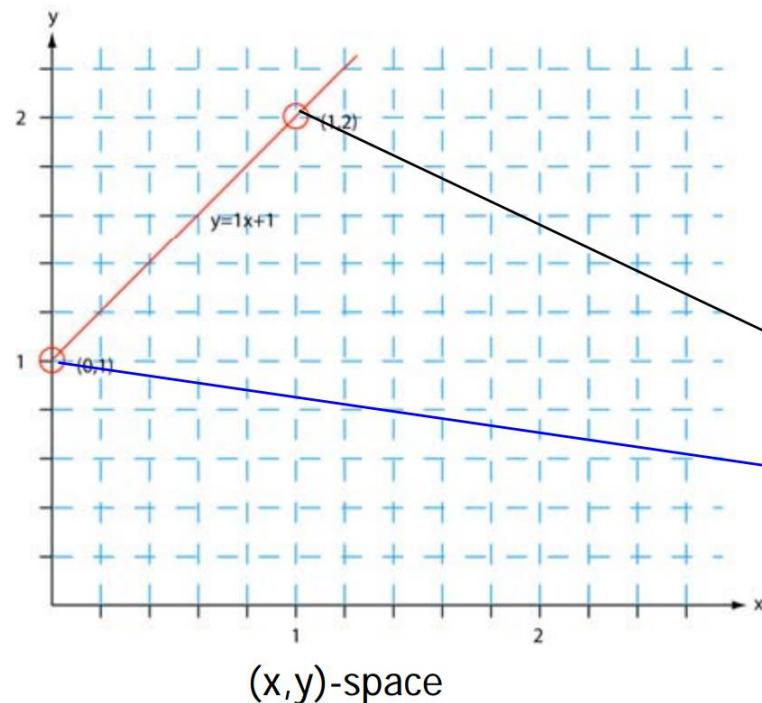
This equation can obviously be rewritten as follows:

- $b = -a^*x_i + y_i$
- We can now consider x and y as parameters
- a and b as variables.

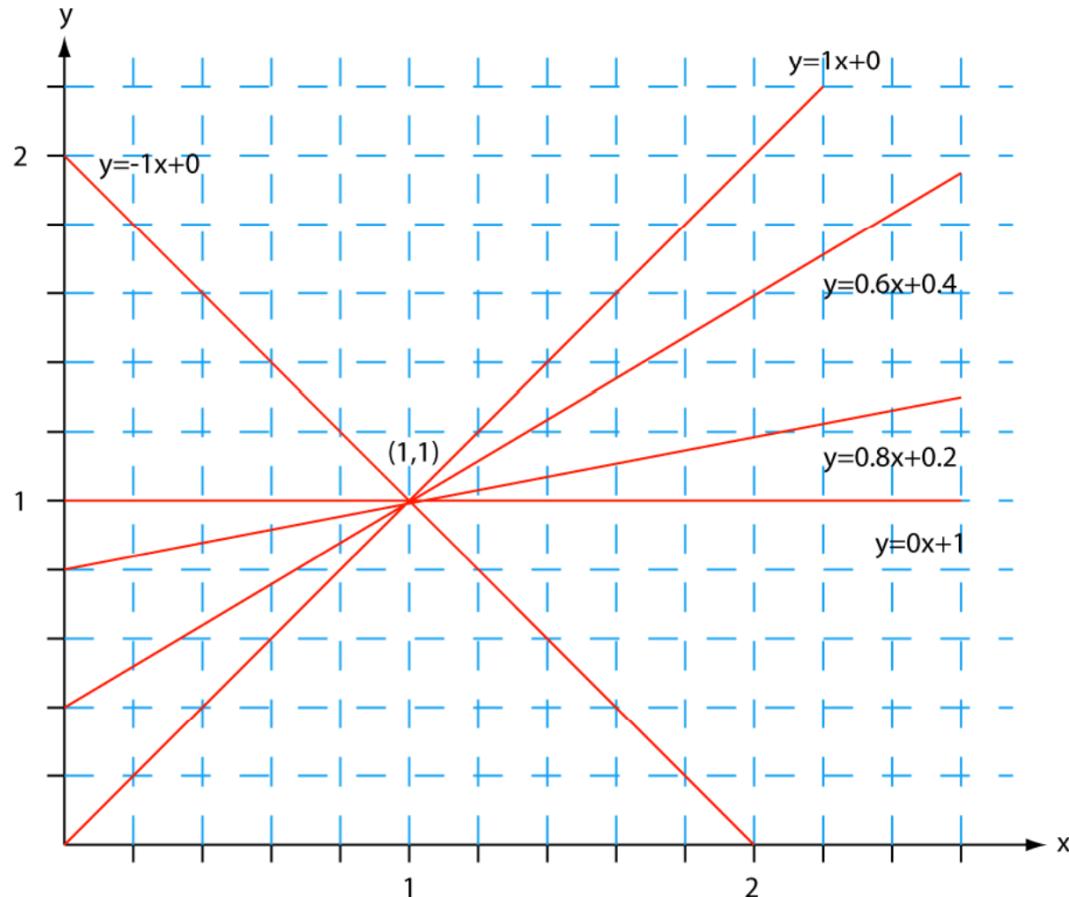
This is a line in (a, b) space parameterized by x and y .

- So: a single point in x_1, y_1 -space gives a line in (a, b) space.
- Another point (x_2, y_2) will give rise to another line (a, b) space.

DETECTING LINES USING HOUGH TRANSFORM



DETECTING LINES USING HOUGH TRANSFORM

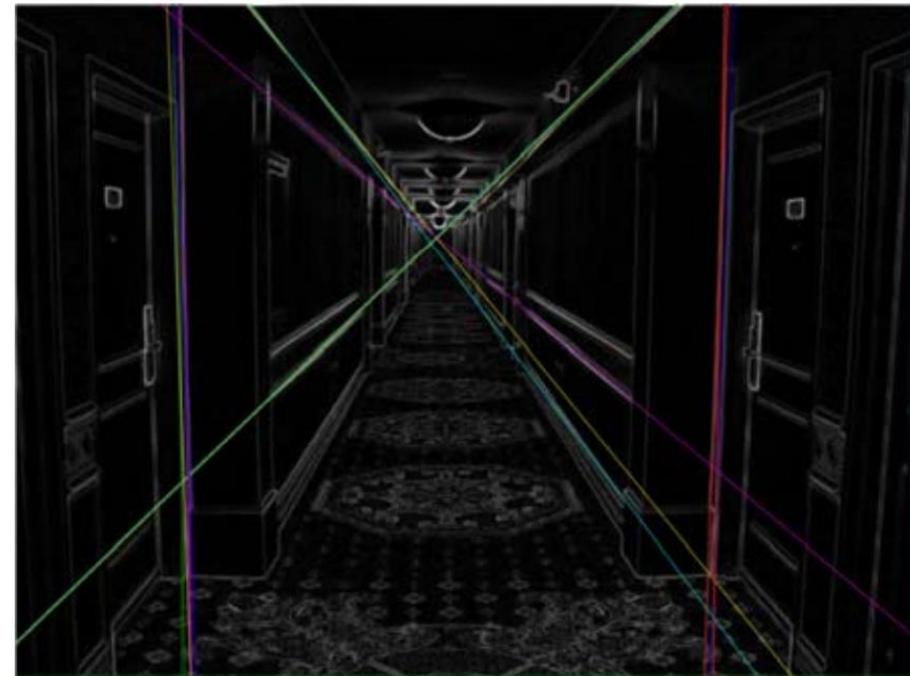


- Two points (x_1, y_1) and (x_2, y_2) define a line in the (x, y) plane.
- These two points give rise to two different lines in (a, b) space.
- In (a, b) space these lines will intersect in a point (a', b')
- All points on the line defined by (x_1, y_1) and (x_2, y_2) in (x, y) space will parameterize lines that intersect in (a', b') in (a, b) space.

1. Quantize the parameter space (a, b) by dividing it into cells
2. This quantized space is often referred to as the accumulator cells.
3. Count the number of times a line intersects a given cell.
 - For each pair of points (x_1, y_1) and (x_2, y_2) detected as an edge, find the intersection (a', b') in (a, b) space.
 - Increase the value of a cell in the range $[[a_{\min}, a_{\max}], [b_{\min}, b_{\max}]]$ that (a', b') belongs to.
 - Cells receiving more than a certain number of counts (also called ‘votes’) are assumed to correspond to lines in (x, y) space.

OUTPUT OF HOUGH TRANSFORM

1. Here are the top 20 most voted lines in the image:



Advantages:

- Conceptually simple.
- Easy implementation
- Handles missing and occluded data very gracefully.
- Can be adapted to many types of forms, not just lines

Disadvantages:

- Computationally complex for objects with many parameters.
- Looks for only one single type of object
- Can be “fooled” by “apparent lines”.
- The length and the position of a line segment cannot be determined.
- Co-linear line segments cannot be separated.

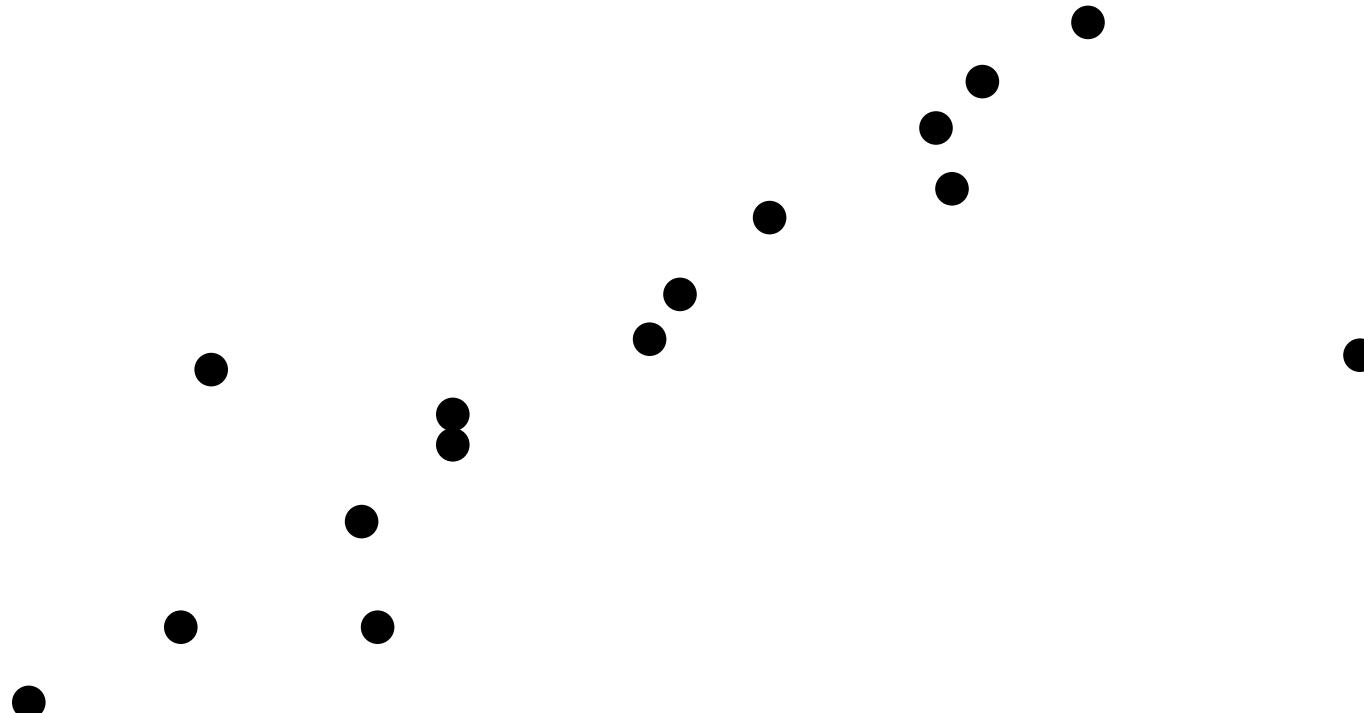
1. RANdom SAmple Consensus
2. Approach: we want to avoid the impact of outliers, so let's look for "inliers", and use only those.
3. Intuition: if an outlier is chosen to compute the current fit, then the resulting line won't have much support from rest of the points.

RANSAC loop:

1. Randomly select a *seed group* of points
2. Compute transformation from seed group
3. Find *inliers* to this transformation
4. If the number of inliers is sufficiently large, re-compute least-squares estimate of transformation on all of the inliers
5. Keep the transformation with the largest number of inliers

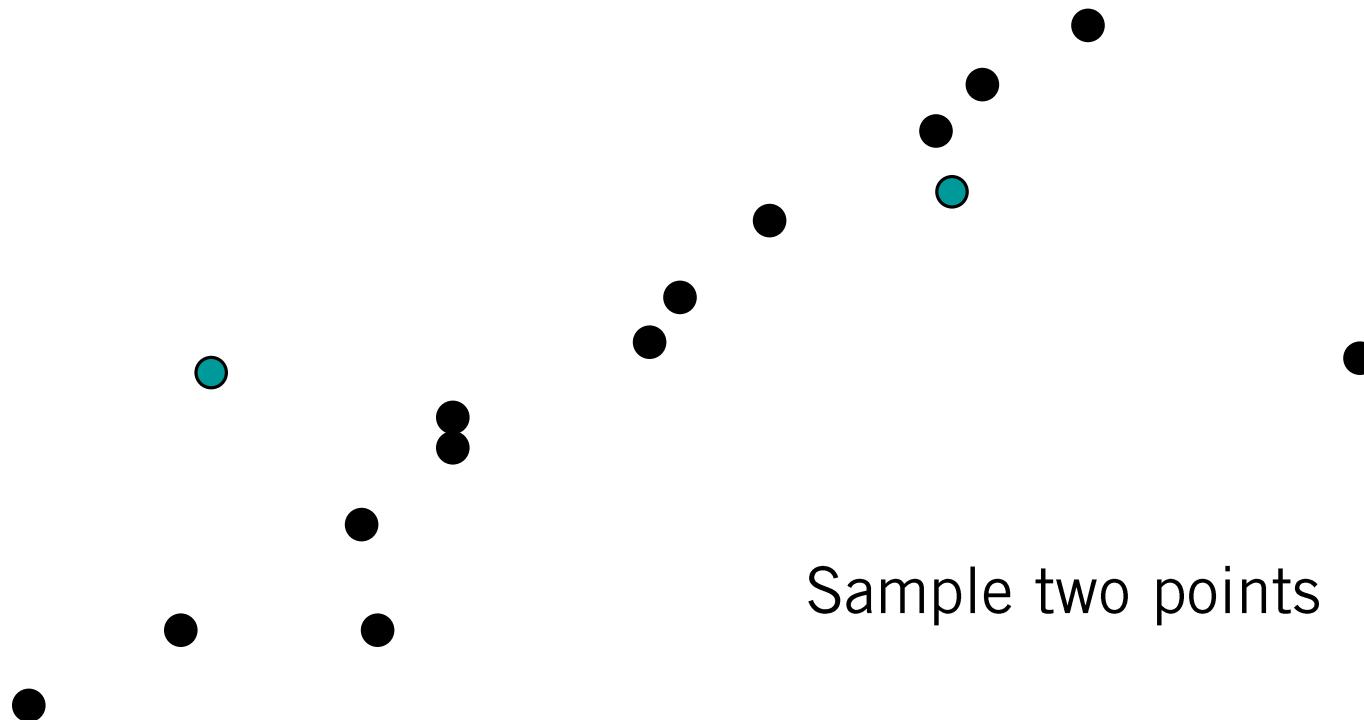
1. Task: Estimate the best line

- *How many points do we need to estimate the line?*



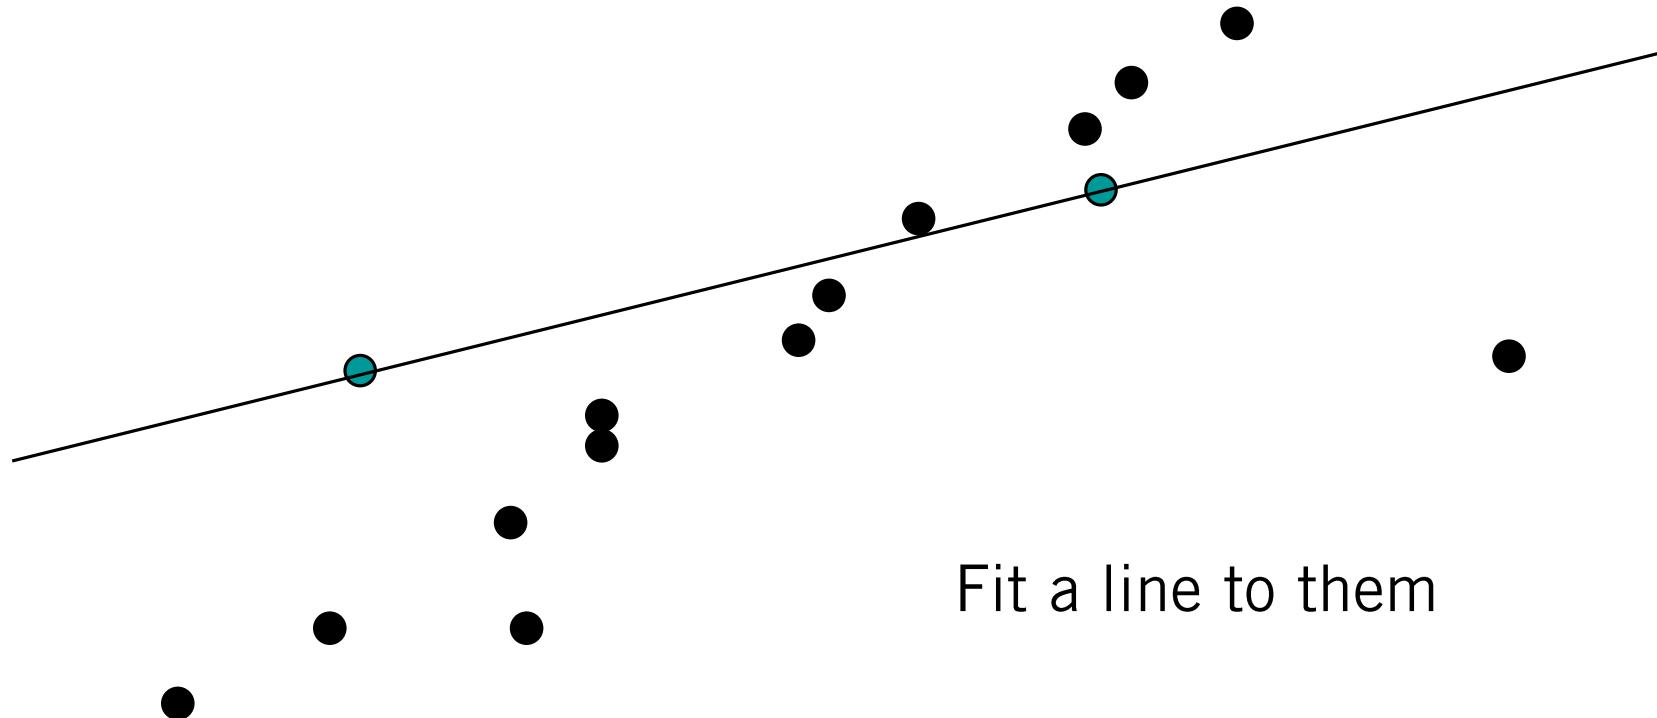
RANSAC LINE FITTING EXAMPLE

1. Task: Estimate the best line

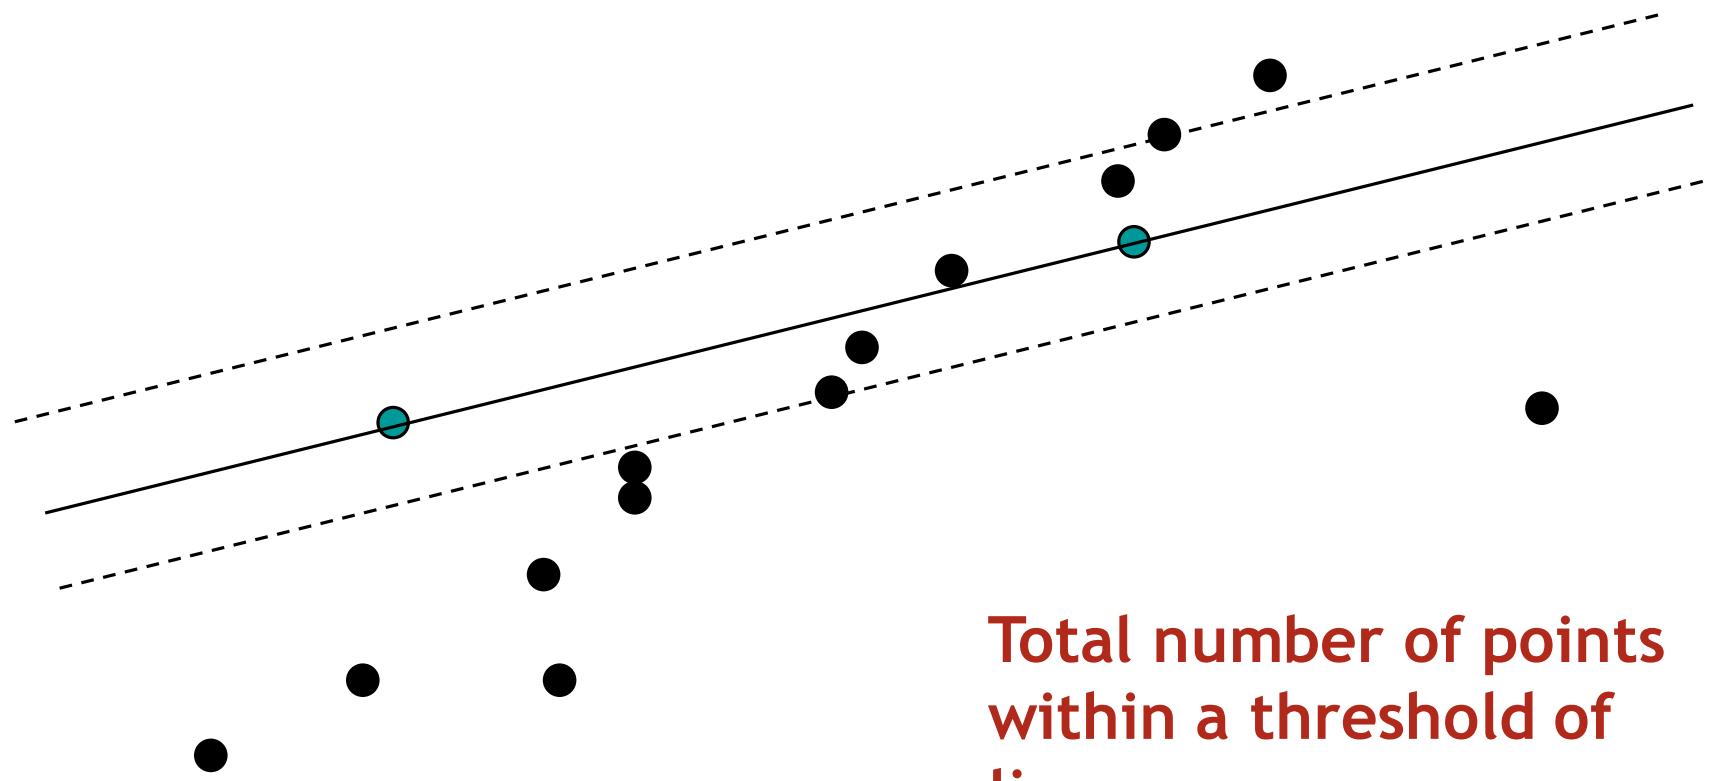


Sample two points

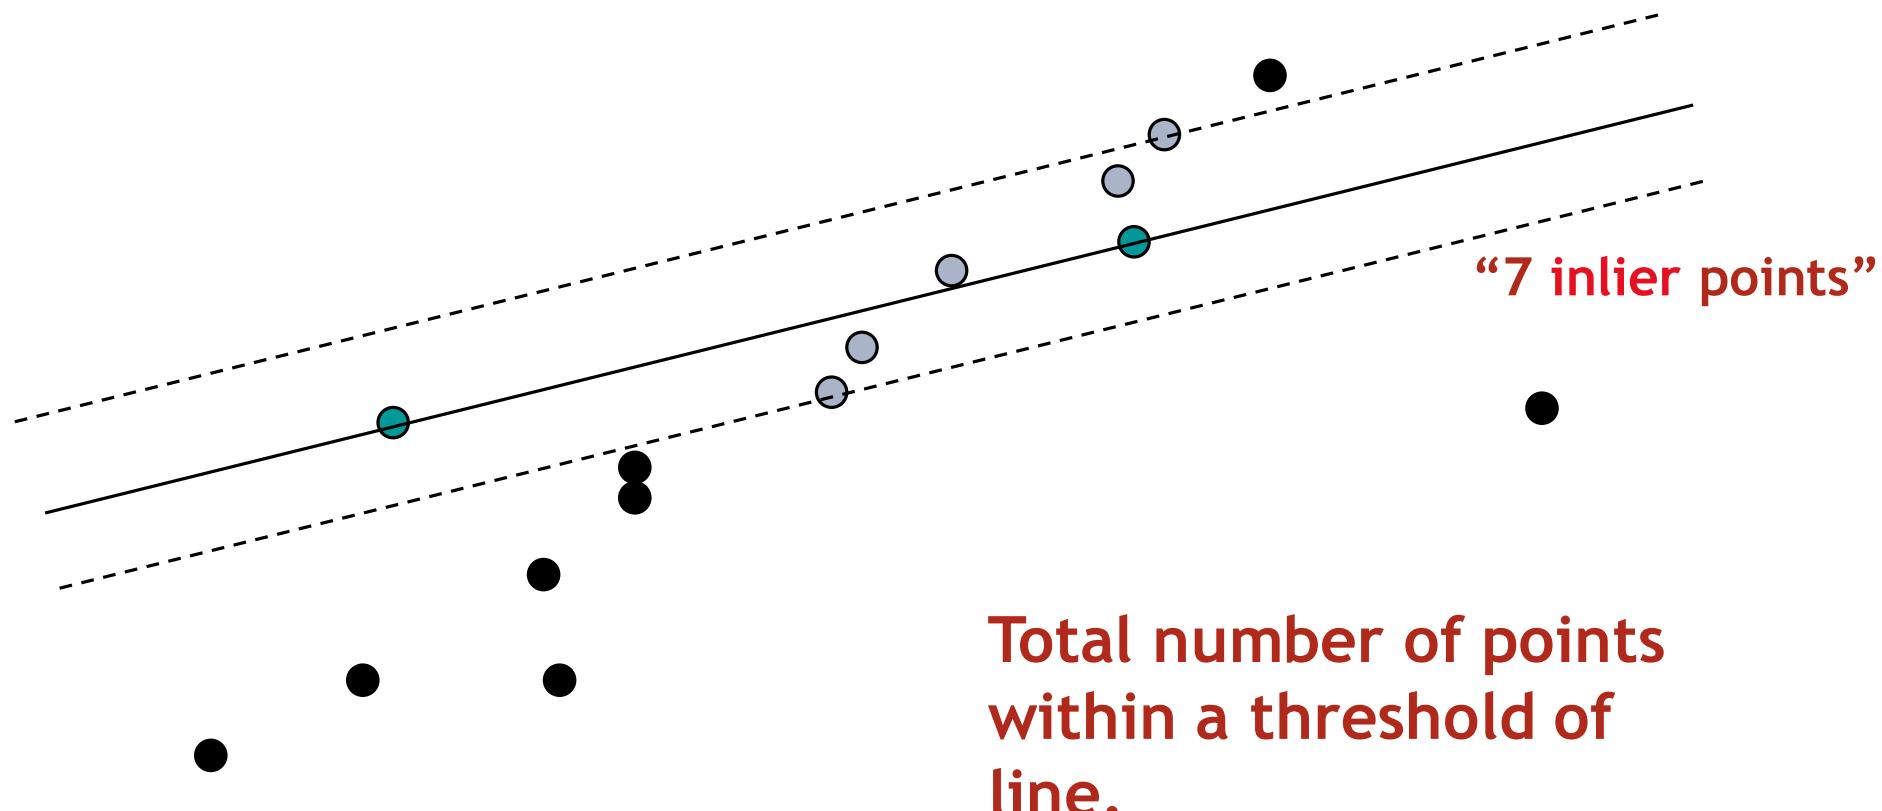
1. Task: Estimate the best line



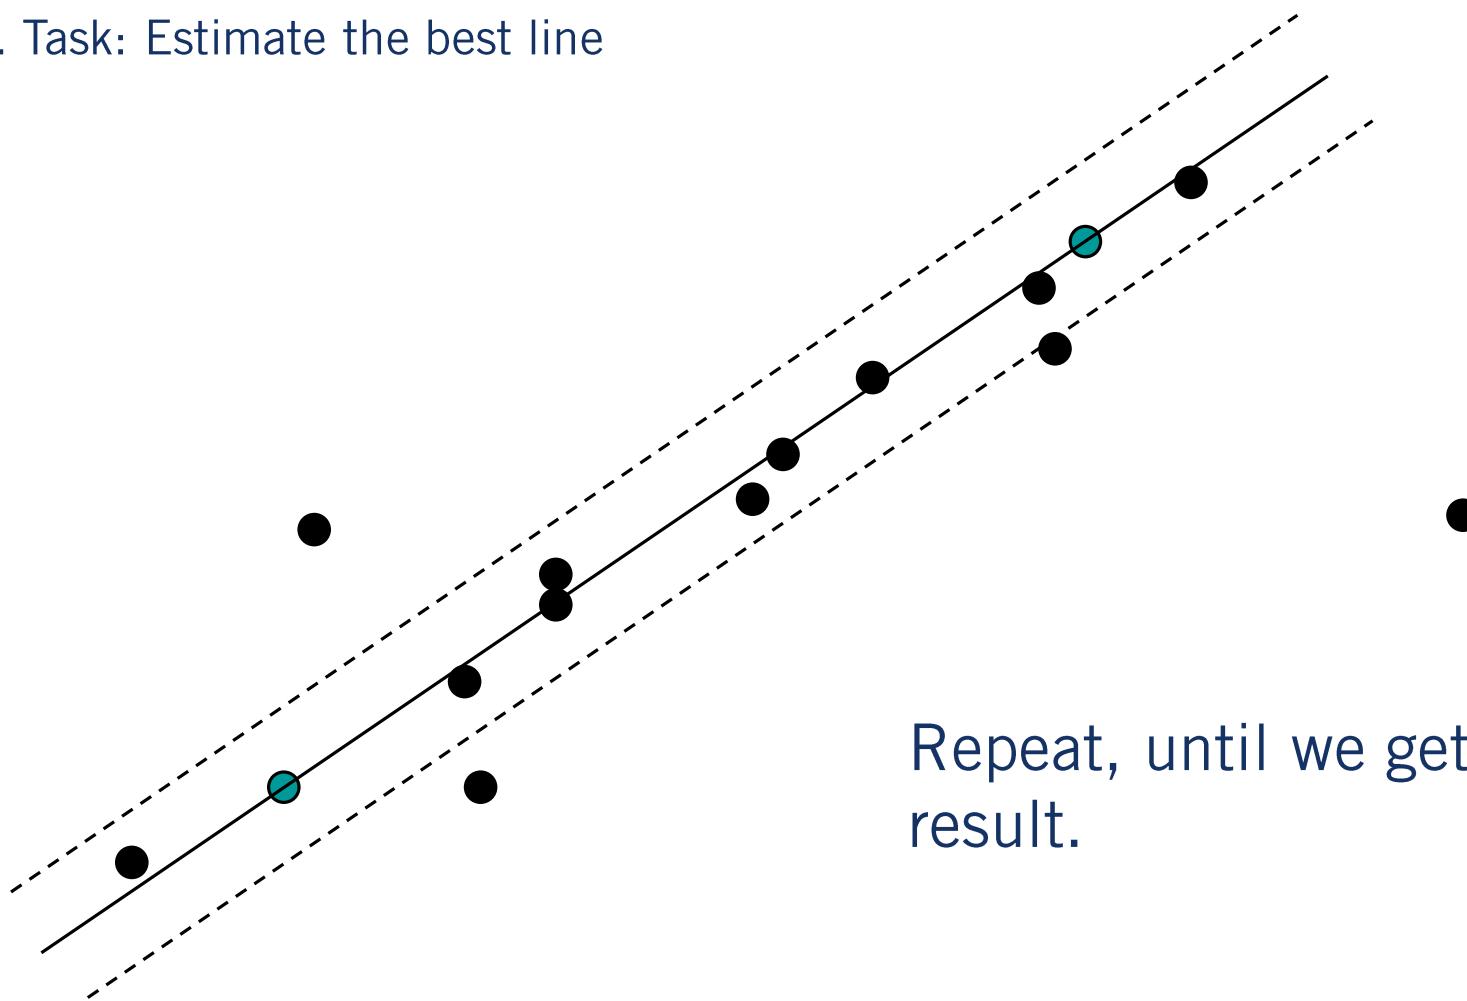
1. Task: Estimate the best line



1. Task: Estimate the best line

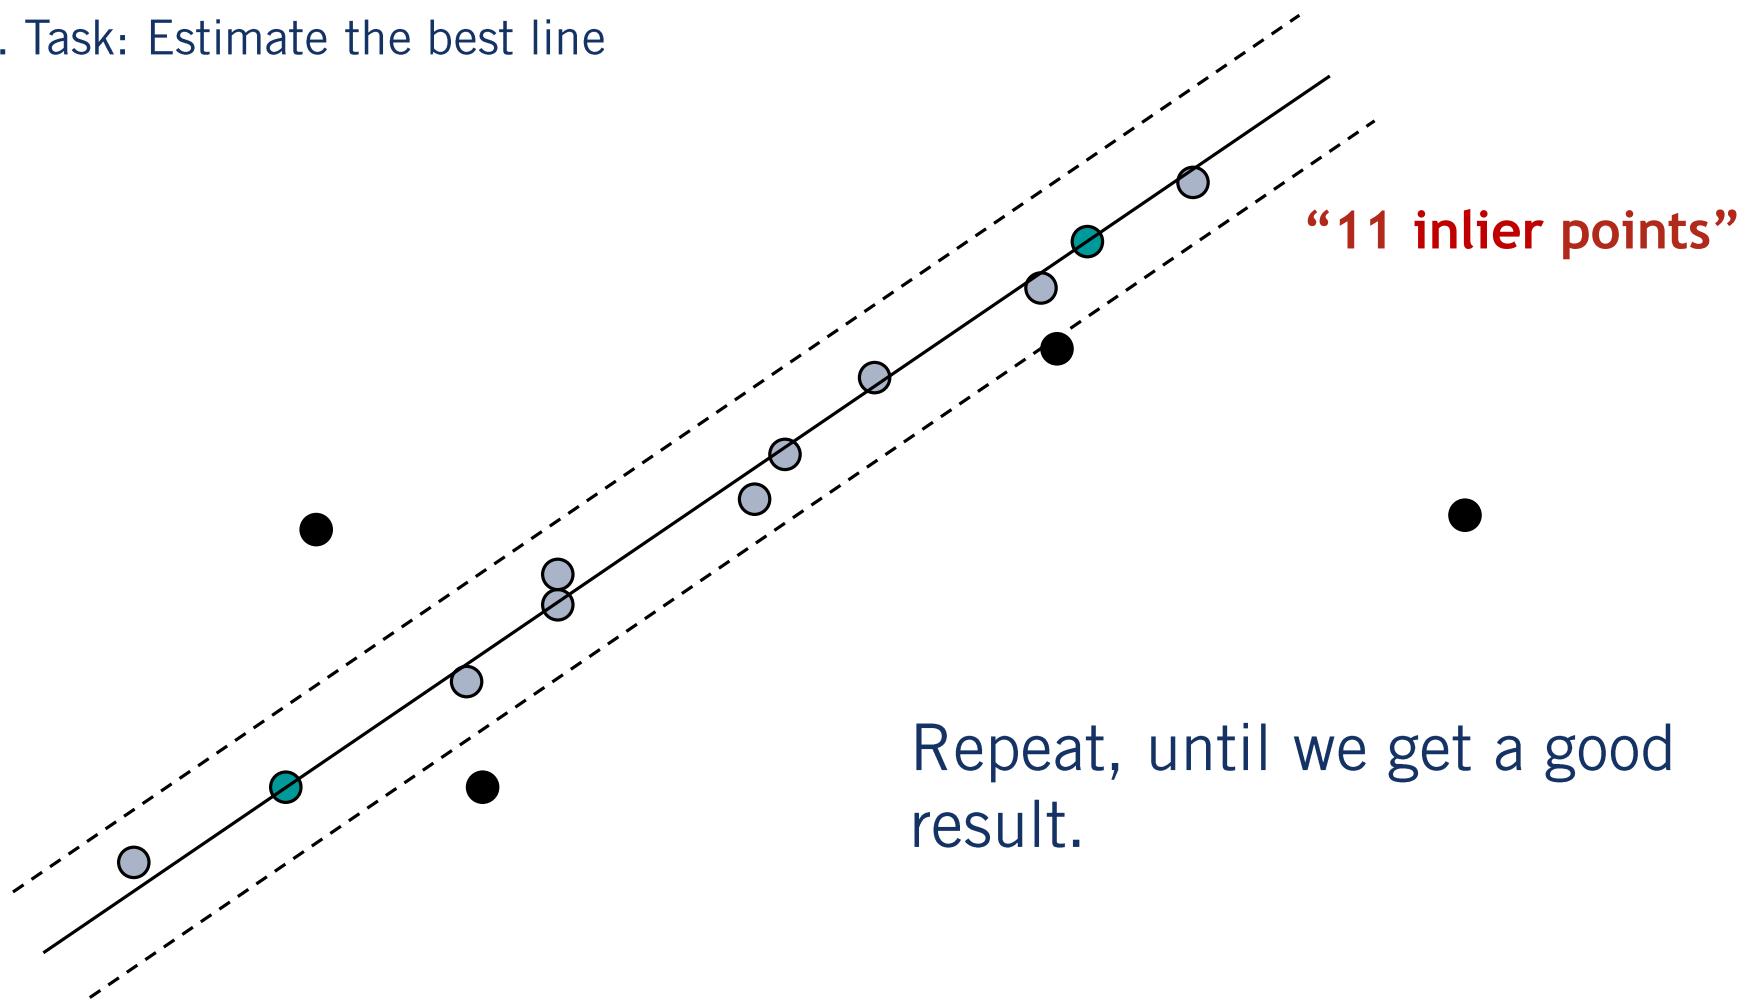


1. Task: Estimate the best line



Repeat, until we get a good result.

1. Task: Estimate the best line



Algorithm 15.4: RANSAC: fitting lines using random sample consensus

Determine:

- n — the smallest number of points required
- k — the number of iterations required
- t — the threshold used to identify a point that fits well
- d — the number of nearby points required
 - to assert a model fits well

Until k iterations have occurred

 Draw a sample of n points from the data

- uniformly and at random

 Fit to that set of n points

 For each data point outside the sample

 Test the distance from the point to the line

- against t ; if the distance from the point to the line
- is less than t , the point is close

 end

 If there are d or more points close to the line

 then there is a good fit. Refit the line using all

- these points.

end

Use the best fit from this collection, using the

- fitting error as a criterion

How many samples are needed?

- Suppose w is fraction of inliers (points from line).
- n points needed to define hypothesis (2 for lines)
- k samples chosen.

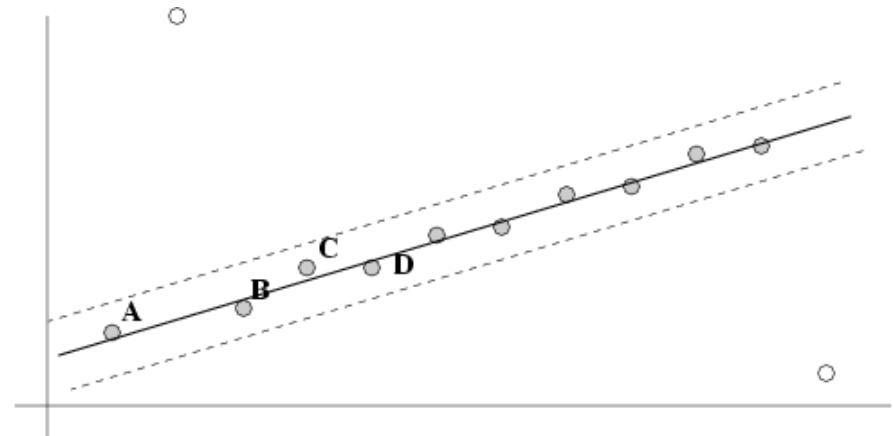
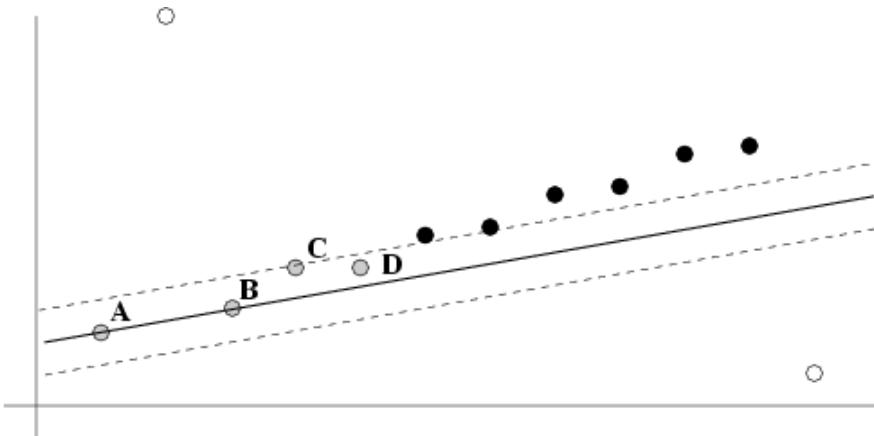
Prob. that a single sample of n points is correct: w^n

1. Prob. that all k samples fail is: $(1-w^n)^k$

⇒ Choose k high enough to keep this below desired failure rate.

Sample size <i>n</i>	Proportion of outliers						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

- 1.RANSAC divides data into inliers and outliers and yields estimate computed from minimal set of inliers.
- 2.Improve this initial estimate with estimation over all inliers (e.g. with standard least-squares minimization).
- 3.But this may change inliers, so alternate fitting with re-classification as inlier/outlier.



- **Pros:**
 - General method suited for a wide range of model fitting problems
 - Easy to implement and easy to calculate its failure rate
- **Cons:**
 - Only handles a moderate percentage of outliers without cost blowing up
 - Many real problems have high rate of outliers (but sometimes selective choice of random subsets can help)
- A voting strategy, The Hough transform, can handle high percentage of outliers

IMAGE MATCHING: A CHALLENGING PROBLEM



IMAGE MATCHING: A CHALLENGING PROBLEM



by Diva Sian

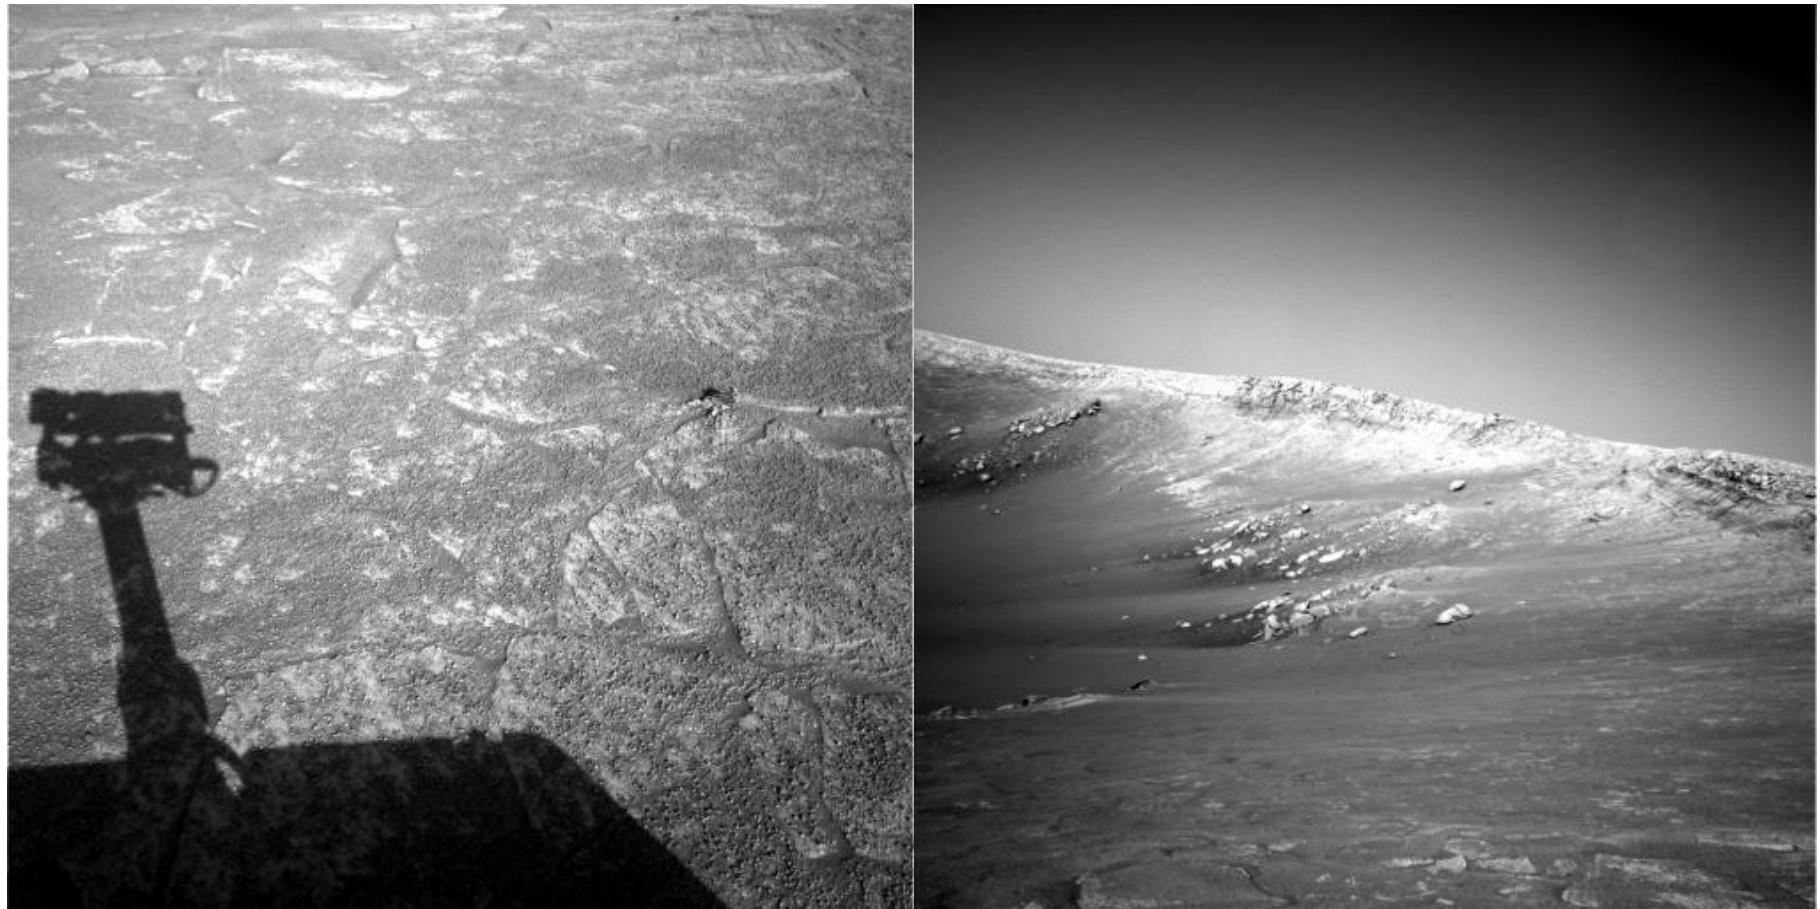


by swashford

HARDER CASE

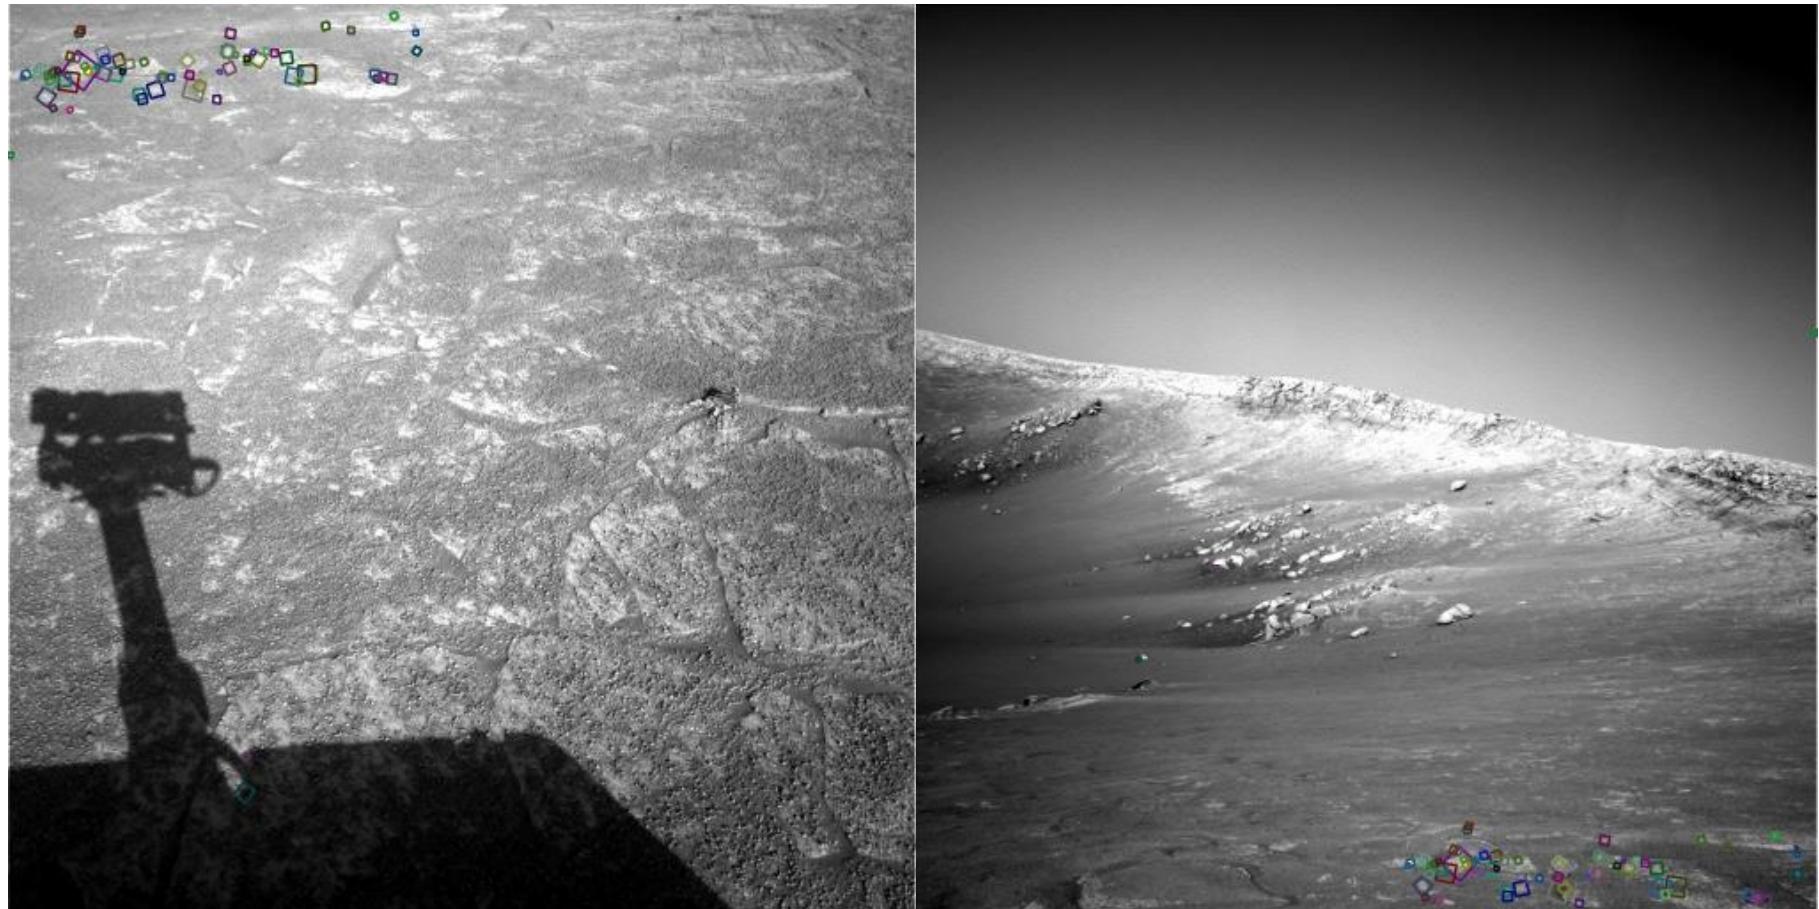


EVEN HARDER ...



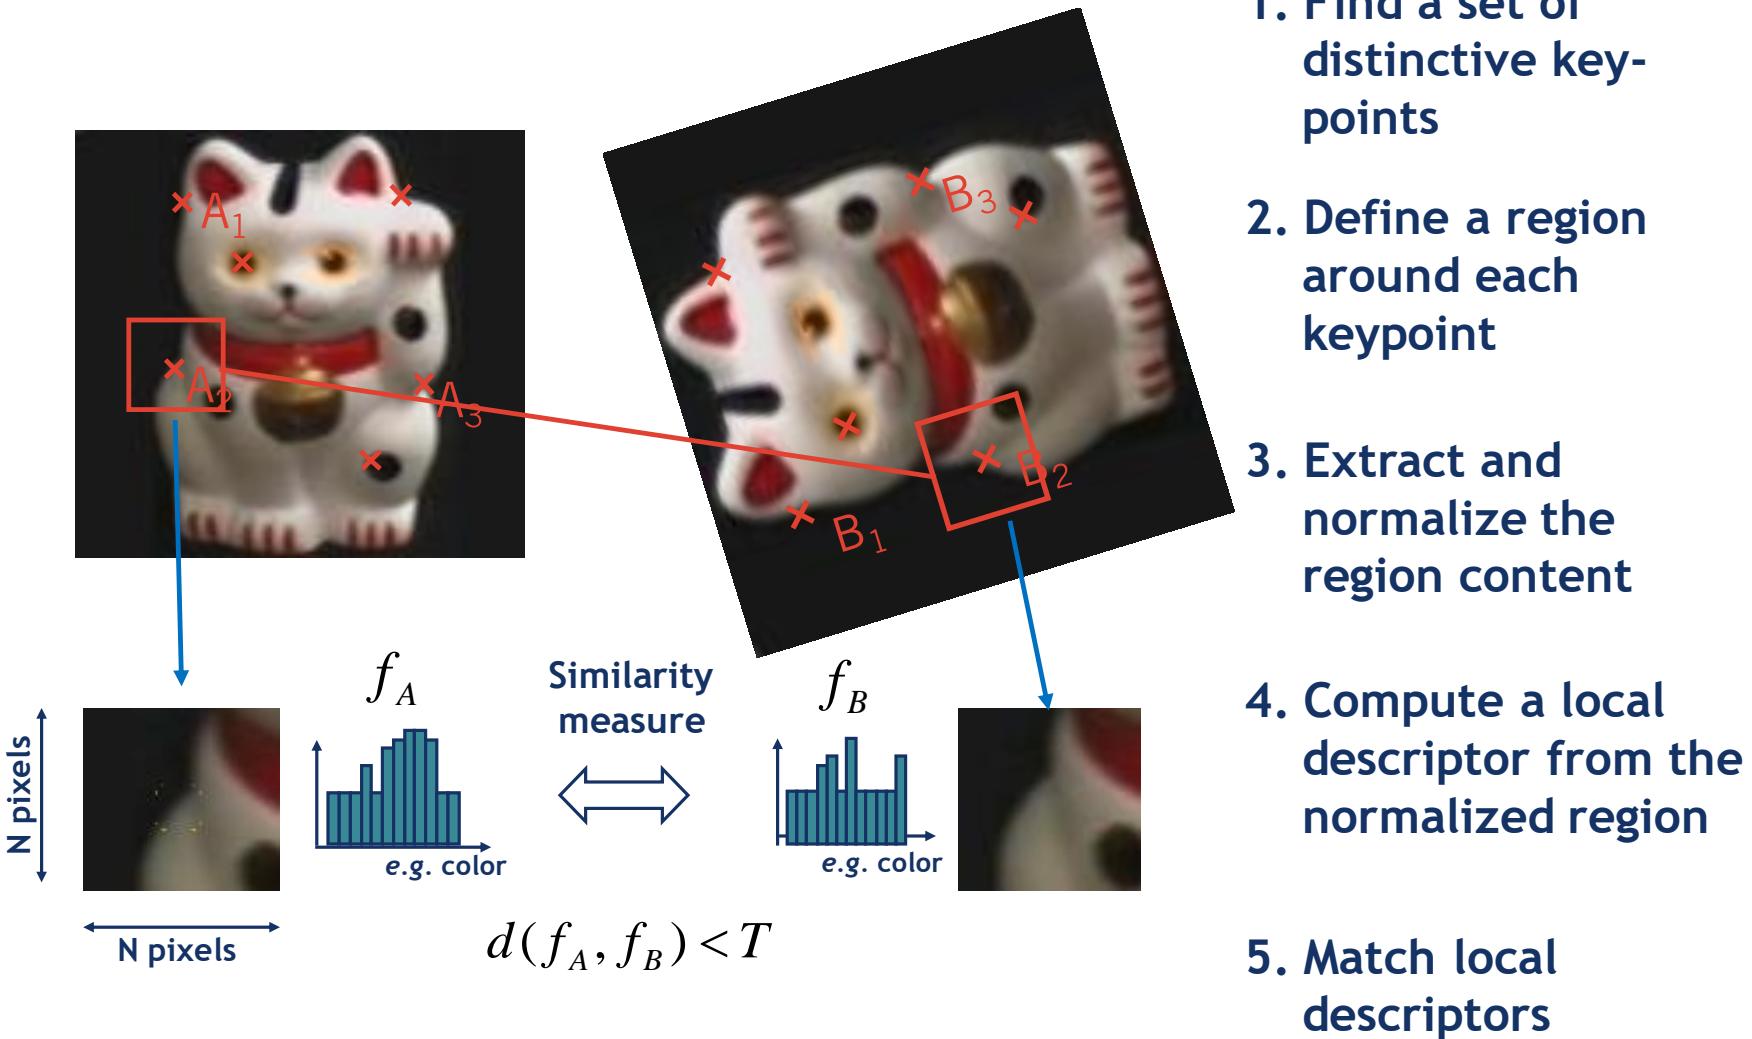
NASA Mars Rover images

EVEN HARDER ...



NASA Mars Rover images with SIFT feature matches
(Figure by Noah Snavely)

GENERAL APPROACH



Problem 1:

- Detect the same point *independently* in both images



No chance to match!

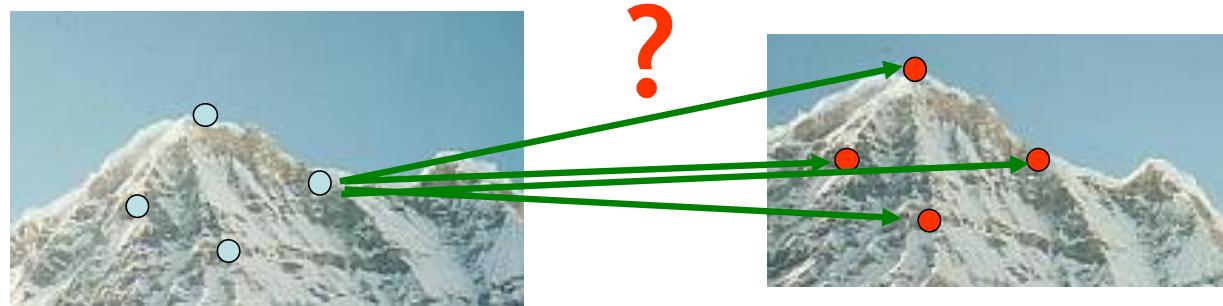
We need a repeatable detector!

Problem 1:

- Detect the same point *independently* in both images

Problem 2:

- For each point correctly recognize the corresponding one



We need a reliable and distinctive descriptor!

INVARIANCE: GEOMETRIC TRANSFORMATIONS

