



# COMPUTER VISION LECTURE 10 – CLUSTERING AND SEGMENTATION

Prof. Dr. Francesco Maurelli  
2018-10-05




# What we will learn today

- Introduction to segmentation and clustering
- Gestalt theory for perceptual grouping
- Agglomerative clustering
- Oversegmentation

**Reading:** [Forsyth & Ponce] Chapters: 14.2, 14.4

# Image Segmentation

- Goal: identify groups of pixels that go together



Slide credit: Steve Seitz, Kristen Grauman





# The Goals of Segmentation

- Separate image into coherent “objects”

Image



Human segmentation





Slide credit: Svetlana Lazebnik

# The Goals of Segmentation

- Separate image into coherent “objects”
- Group together similar-looking pixels for efficiency of further processing


“superpixels”




X. Ren and J. Malik. [Learning a classification model for segmentation](#). ICCV 2003.

Slide credit: Svetlana Lazebnik

# Segmentation for efficiency



[Felzenszwalb and Huttenlocher 2004]



[Shi and Malik 2001]

[Hoiem et al. 2005, Mori 2005]

Slide: Derek Hoiem



# Segmentation as a result




Rother et al. 2004




# Types of segmentations



Oversegmentation



Undersegmentation



Multiple Segmentations



# One way to think about “segmentation” is Clustering

Clustering: group together similar data points and represent them with a single token

Key Challenges:

- 1) What makes two points/images/patches similar?
- 2) How do we compute an overall grouping from pairwise similarities?

Slide: Derek Hoiem

# Why do we cluster?

- **Summarizing data**
  - Look at large amounts of data
  - Patch-based compression or denoising
  - Represent a large continuous vector with the cluster number
- **Counting**
  - Histograms of texture, color, SIFT vectors
- **Segmentation**
  - Separate the image into different regions
- **Prediction**
  - Images in the same cluster may have the same labels

Slide: Derek Hoiem

# How do we cluster?


- Agglomerative clustering
  - Start with each point as its own cluster and iteratively merge the closest clusters
- K-means (next lecture)
  - Iteratively re-assign points to the nearest cluster center
- Mean-shift clustering (next lecture)
  - Estimate modes of pdf



# General ideas

- Tokens
    - whatever we need to group (pixels, points, surface elements, etc., etc.)
  - Bottom up clustering
    - tokens belong together because they are locally coherent
  - Top down clustering
    - tokens belong together because they lie on the same visual entity (object, scene...)
- > These two are not mutually exclusive


# Examples of Grouping in Vision



## Determining image regions




## Grouping video frames into shots



## Figure-ground

## What things should be grouped?

## What cues indicate groups?



## Object-level grouping

Slide credit: Kristen Grauman



# Similarity



Slide credit: Kristen Grauman



# Symmetry



Slide credit: Kristen Grauman

# Common Fate



Image credit: Arthus-Bertrand (via F. Durand)



(c) 2006 Helko Burkhart, illiano.com

Slide credit: Kristen Grauman




# Proximity



Slide credit: Kristen Grauman



# Muller-Lyer Illusion




- What makes the bottom line look longer than the top line?

# What we will learn today

- Introduction to segmentation and clustering
- **Gestalt theory for perceptual grouping**
- Agglomerative clustering
- Oversegmentation

# The Gestalt School

- Grouping is key to visual perception
- Elements in a collection can have properties that result from **relationships**
  - “The whole is greater than the sum of its parts”




[http://en.wikipedia.org/wiki/Gestalt\\_psychology](http://en.wikipedia.org/wiki/Gestalt_psychology)

# Gestalt Theory

- Gestalt: whole or group
  - Whole is greater than sum of its parts
  - Relationships among parts can yield new properties/features
- Psychologists identified series of factors that predispose set of elements to be grouped (by human visual system)

*“I stand at the window and see a house, trees, sky.  
Theoretically I might say there were 327 brightnesses  
and nuances of colour. Do I have “327”? No. I have sky, house,  
and trees.”*


**Max Wertheimer**  
(1880-1943)



Untersuchungen zur Lehre von der Gestalt,  
*Psychologische Forschung*, Vol. 4, pp. 301-350, 1923  
<http://psy.ed.asu.edu/~classics/Wertheimer/Forms/forms.htm>




# Gestalt Factors



Not grouped



Proximity




Similarity




Similarity




Common Fate




Common Region




Parallelism



Symmetry




Continuity




Closure

- These factors make intuitive sense, but are very difficult to translate into algorithms.

# Continuity through Occlusion Cues




# Continuity through Occlusion Cues



Continuity, explanation by occlusion

# Figure-Ground Discrimination





# The Ultimate Gestalt?



# What we will learn today

- Introduction to segmentation and clustering
- Gestalt theory for perceptual grouping
- **Agglomerative clustering**
- Oversegmentation

A composite image showing a woman on the left and a dog on the right. Both have long, curly, light brown hair. The woman is looking directly at the camera with a neutral expression. The dog is also looking directly at the camera. The background is a solid, muted brown color.

# What is similarity?

Similarity is hard to define, but... “We know it when we see it” The real meaning of similarity is a philosophical question. We will take a more pragmatic approach.

# Clustering: distance measure


Clustering is an unsupervised learning method. Given items  $x_1, \dots, x_n \in \mathbb{R}^D$ , the goal is to group them into clusters. We need a pairwise distance/similarity function between items, and sometimes the desired number of clusters.



# Desirable Properties of a Clustering Algorithms

- Scalability (in terms of both time and space)
- Ability to deal with different data types
- Minimal requirements for domain knowledge to determine input parameters
- Interpretability and usability Optional
  - Incorporation of user-specified constraints

# Animated example



[source](#)

# Animated example
























Diagram illustrating a sequence of characters and their relationships. The characters are arranged in a sequence, with lines indicating connections between them. The characters are: a woman in a green dress, a young girl in a red dress, a woman in a green dress, a woman in a blue dress, and a woman in a pink dress.

				
0	8	8	7	7
	0	2	4	4
		0	3	3
			0	1
				0
				

[source](#)

# Animated example




				
0	8	8	7	7
	0	2	4	4
		0	3	3
			0	1
				0

[source](#)




# Agglomerative clustering



1. Say "Every point is its own cluster"

Slide credit: Andrew Moore

# Agglomerative clustering




1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters



Slide credit: Andrew Moore

# Agglomerative clustering




1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters
3. Merge it into a parent cluster



Slide credit: Andrew Moore

# Agglomerative clustering




1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters
3. Merge it into a parent cluster
4. Repeat



Slide credit: Andrew Moore



# Agglomerative clustering



1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters
3. Merge it into a parent cluster
4. Repeat




Slide credit: Andrew Moore

# Agglomerative clustering


## How to define cluster similarity?

- Average distance between points,
- maximum distance
- minimum distance
- Distance between means or medoids



## How many clusters?

- Clustering creates a dendrogram (a tree)
- Threshold based on max number of clusters or based on distance between merges




# Agglomerative Hierarchical Clustering - Algorithm

1. Initially each item  $x_1, \dots, x_n$  is in its own cluster  $C_1, \dots, C_n$ .
2. Repeat until there is only one cluster left:
3.       Merge the nearest clusters, say  $C_i$  and  $C_j$ .

# Different measures of nearest clusters

## Single Link

- $d(C_i, C_j) = \min_{x \in C_i, x' \in C_j} d(x, x')$ . This is known as *single-linkage*. It is equivalent to the minimum spanning tree algorithm. One can set a threshold and stop clustering once the distance between clusters is above the threshold. Single-linkage tends to produce long and skinny clusters.




Long, skinny clusters



# Different measures of nearest clusters

## Complete Link

- $d(C_i, C_j) = \max_{x \in C_i, x' \in C_j} d(x, x')$ . This is known as *complete-linkage*. Clusters tend to be compact and roughly equal in diameter.




Tight clusters

# Different measures of nearest clusters

## Average Link

- $d(C_i, C_j) = \frac{\sum_{x \in C_i, x' \in C_j} d(x, x')}{|C_i| \cdot |C_j|}$ . This is the average distance between items. Somewhere between single-linkage and complete-linkage.



Robust against noise.

# Conclusions: Agglomerative Clustering

## Good

- Simple to implement, widespread application.
- Clusters have adaptive shapes.
- Provides a hierarchy of clusters.
- No need to specify number of clusters in advance.

## Bad

- May have imbalanced clusters.
- Still have to choose number of clusters or threshold.
- Does not scale well. Runtime of  $O(n^3)$ .
- Can get stuck at a local optima.

# What we will learn today?

- Introduction to segmentation and clustering
- Gestalt theory for perceptual grouping
- Agglomerative clustering
- **Oversegmentation**


# How do we segment using Clustering?

- Solution: Oversegmentation algorithm
  - Introduced by *Felzenszwalb and Huttenlocher* in the paper titled *Efficient Graph-Based Image Segmentation*.






# Problem Formulation



- Graph  $G = (V, E)$
- $V$  is set of nodes (i.e. pixels)
- $E$  is a set of undirected edges between pairs of pixels
- $w(v_i, v_j)$  is the weight of the edge between nodes  $v_i$  and  $v_j$ .
- $S$  is a segmentation of a graph  $G$  such that  $G' = (V, E')$  where  $E' \subset E$ .
- $S$  divides  $G$  into  $G'$  such that it contains distinct clusters  $C$ .

# Predicate for Segmentation




- Predicate D determines whether there is a boundary for segmentation.

$$Merge(C_1, C_2) = \begin{cases} True & \text{if } dif(C_1, C_2) < in(C_1, C_2) \\ False & \text{otherwise} \end{cases}$$

Where

- $dif(C_1, C_2)$  is the difference between two clusters.
- $in(C_1, C_2)$  is the internal different in the clusters  $C_1$  and  $C_2$

# Predicate for Segmentation




- Predicate D determines whether there is a boundary for segmentation.

$$Merge(C_1, C_2) = \begin{cases} True & \text{if } dif(C_1, C_2) < in(C_1, C_2) \\ False & \text{otherwise} \end{cases}$$

$$dif(C_1, C_2) = \min_{v_i \in C_1, v_j \in C_2, (C_1, C_2) \in E} w(v_i, v_j)$$

The different between two components is the minimum weight edge that connects a node  $v_i$  in clusters  $C_1$  to node  $v_j$  in  $C_2$

# Predicate for Segmentation



- Predicate D determines whether there is a boundary for segmentation.

$$Merge(C_1, C_2) = \begin{cases} True & \text{if } dif(C_1, C_2) < in(C_1, C_2) \\ False & \text{otherwise} \end{cases}$$



$$dif(C_1, C_2) = \min_{v_i \in C_1, v_j \in C_2, (C_1, C_2) \in E} w(v_i, v_j)$$

$$in(C_1, C_2) = \min_{C \in \{C_1, C_2\}} \left[ \max_{v_i, v_j \in C} \left[ w(v_i, v_j) + \frac{k}{|C|} \right] \right]$$

$in(C_1, C_2)$  is to the maximum weight edge that connects two nodes in the same component.

# Predicate for Segmentation


- $k/|C|$  sets the threshold by which the components need to be different from the internal nodes in a component.
- Properties of constant  $k$ :
  - If  $k$  is large, it causes a preference of larger objects.
  - $k$  does not set a minimum size for components.







# Features and weights

- Project every pixel into feature space defined by  $(x, y, r, g, b)$ .
- Every pixel is connected to its 8 neighboring pixels and the weights are determined by the difference in intensities.
- Weights between pixels are determined using L2 (Euclidian) distance in feature space.
- Edges are chosen for only top ten nearest neighbors in feature space to ensure run time of  $O(n \log n)$  where  $n$  is number of pixels.



# Results



# What we have learned today?

- Introduction to segmentation and clustering
- Gestalt theory for perceptual grouping
- Agglomerative clustering
- Oversegmentation