

COMPUTER VISION LECTURE 21 – FACE RECOGNITION

Prof. Dr. Francesco Maurelli
2018-11-20




What we will learn today

- Introduction to face recognition
- The Eigenfaces Algorithm
- Linear Discriminant Analysis (LDA)


Turk and Pentland, Eigenfaces for Recognition, *Journal of Cognitive Neuroscience* **3** (1): 71–86.

P. Belhumeur, J. Hespanha, and D. Kriegman. "Eigenfaces vs. Fisherfaces: Recognition Using Class Specific Linear Projection". *IEEE Transactions on pattern analysis and machine intelligence* **19** (7): 711. 1997.

Detection versus Recognition



Detection finds the faces in images



Recognition recognizes WHO the person is

Face Recognition

- Digital photography



Face Recognition

- Digital photography
- Surveillance



The interface displays a video recording of two individuals walking down a corridor. Red bounding boxes are drawn around their faces, indicating detection. Below the video, a red square icon is followed by the word "Recording".

At the bottom left, there is a button labeled "Report".

On the right side, a section titled "Detecting...." shows two small portrait images of the detected faces. Below this, a section titled "Matching with Database" displays two results:

- The first result shows a portrait of a man with the following details:
 - Name: Alireza,
 - Date: 25 My 2007 15:45
 - Place: Main corridor
- The second result shows a portrait of a man with the following details:
 - Name: **Unknown**
 - Date: 25 My 2007 15:45
 - Place: Main corridor


Face Recognition

- Digital photography
- Surveillance
- Album organization




Face Recognition

- Digital photography
- Surveillance
- Album organization
- Person tracking/id.



Face Recognition

- Digital photography
- Surveillance
- Album organization
- Person tracking/id.
- Emotions and expressions




Face Recognition

- Digital photography
- Surveillance
- Album organization
- Person tracking/id.
- Emotions and expressions
- Security/warfare
- Tele-conferencing
- Etc.


The Space of Faces

- An image is a point in a high dimensional space
 - If represented in grayscale intensity, an $N \times M$ image is a point in \mathbb{R}^{NM}
 - E.g. 100x100 image = 10,000 dim




Slide credit: Chuck Dyer, Steve Seitz, Nishino

100x100 images can contain many things other than faces!




The Space of Faces



- An image is a point in a high dimensional space
 - If represented in grayscale intensity, an $N \times M$ image is a point in R^{NM}
 - E.g. 100x100 image = 10,000 dim
- However, relatively few high dimensional vectors correspond to valid face images
- We want to effectively model the subspace of face images

Slide credit: Chuck Dyer, Steve Seitz, Nishino

Where have we seen something like this before?



- A face image
- A (non-face) image

Image
space

Face space



- Compute n-dim subspace such that the projection of the data points onto the subspace has **the largest variance** among all n-dim subspaces.
- **Maximize the scatter** of the training images in face space

Key Idea

- So, compress them to a low-dimensional subspace that captures key appearance characteristics of the visual DOFs.
- **USE PCA for estimating the sub-space**
(dimensionality reduction)
- Compare two faces by projecting the images into the subspace and measuring the EUCLIDEAN distance between them.

What we will learn today

- Introduction to face recognition
- The Eigenfaces Algorithm
- Linear Discriminant Analysis (LDA)

Turk and Pentland, Eigenfaces for Recognition, *Journal of Cognitive Neuroscience* **3** (1): 71–86.

P. Belhumeur, J. Hespanha, and D. Kriegman. "Eigenfaces vs. Fisherfaces: Recognition Using Class Specific Linear Projection". *IEEE Transactions on pattern analysis and machine intelligence* **19** (7): 711. 1997.

Eigenfaces: key idea

- Assume that most face images lie on a low-dimensional subspace determined by the first k ($k \ll d$) directions of maximum variance
- Use PCA to determine the vectors or “eigenfaces” that span that subspace
- Represent all face images in the dataset as linear combinations of eigenfaces


M. Turk and A. Pentland, [Face Recognition using Eigenfaces](#), CVPR 1991

Training images: $\mathbf{x}_1, \dots, \mathbf{x}_N$



Top eigenvectors: ϕ_1, \dots, ϕ_k

Mean: μ



Eigenface algorithm

- Training

1. Align training images x_1, x_2, \dots, x_N



Note that each image is formulated into a long vector!

2. Compute average face $m = \frac{1}{N} \sum x_i$
3. Compute the difference image (the centered data matrix)

$$X_c = \begin{bmatrix} | & & | \\ x_1 & \dots & x_n \\ | & & | \end{bmatrix} - \begin{bmatrix} | & & | \\ \mu & \dots & \mu \\ | & & | \end{bmatrix}$$

Eigenface algorithm

4. Compute the covariance matrix

$$\Sigma = \frac{1}{n} \begin{bmatrix} | & & | \\ x_1^c & \dots & x_n^c \\ | & & | \end{bmatrix} \begin{bmatrix} - & x_1^c & - \\ \vdots & & \\ - & x_n^c & - \end{bmatrix} = \frac{1}{n} X_c X_c^T$$

5. Compute the eigenvectors of the covariance matrix Σ
6. Compute each training image x_i 's projections as

$$x_i \rightarrow (x_i^c \cdot f_1, x_i^c \cdot f_2, \dots, x_i^c \cdot f_K) \equiv (a_1, a_2, \dots, a_K)$$

7. Visualize the estimated training face x_i

$$x_i \approx m + a_1 f_1 + a_2 f_2 + \dots + a_K f_K$$

Eigenface algorithm




6. Compute each training image x_i 's projections as

$$x_i \rightarrow (x_i^c \cdot f_1, x_i^c \cdot f_2, \dots, x_i^c \cdot f_K) \equiv (a_1, a_2, \dots, a_K)$$

7. Visualize the reconstructed training face x_i

$$x_i \approx m + a_1 f_1 + a_2 f_2 + \dots + a_K f_K$$

Eigenvalues (variance along eigenvectors)



Reconstruction and Errors

$K = 4$



$K = 200$



$K = 400$



- Only selecting the top K eigenfaces \rightarrow reduces the dimensionality.
- Fewer eigenfaces result in more information loss, and hence less discrimination between faces.

Eigenface algorithm

- Testing

1. Take query image t
2. Project into eigenface space and compute projection

$$t \rightarrow \left((t - m) \cdot f_1, (t - m) \cdot f_2, \dots, (t - m) \cdot f_K \right) \equiv (w_1, w_2, \dots, w_K)$$

3. Compare projection w with all N training projections

- Simple comparison metric: Euclidean
- Simple decision: K-Nearest Neighbor

(note: this “K” refers to the k-NN algorithm, is different from the previous K’s referring to the # of principal components)

Shortcomings

- Requires carefully controlled data:
 - All faces centered in frame
 - Same size
 - Some sensitivity to angle
- Alternative:
 - “Learn” one set of PCA vectors for each angle
 - Use the one with lowest error
- Method is completely knowledge free
 - (sometimes this is good!)
 - Doesn't know that faces are wrapped around 3D objects (heads)

Summary for Eigenface

Pros

- Non-iterative, globally optimal solution

Limitations

- PCA projection is **optimal for reconstruction** from a low dimensional basis, but **may NOT be optimal for discrimination...**

Besides face recognitions,
we can also do
Facial expression recognition

Happiness subspace (method A)



Disgust subspace (method A)



Facial Expression Recognition Movies (method A)




What we will learn today

- Introduction to face recognition
- The Eigenfaces Algorithm
- Linear Discriminant Analysis (LDA)

Turk and Pentland, Eigenfaces for Recognition, *Journal of Cognitive Neuroscience* **3** (1): 71–86.


P. Belhumeur, J. Hespanha, and D. Kriegman. "Eigenfaces vs. Fisherfaces: Recognition Using Class Specific Linear Projection". *IEEE Transactions on pattern analysis and machine intelligence* **19** (7): 711. 1997.

Which direction is the first principle component?



Fischer's Linear Discriminant Analysis

- Goal: find the best separation between two classes




Slide inspired by N. Vasconcelos

Difference between PCA and LDA


- PCA preserves maximum variance
- LDA preserves discrimination
 - Find projection that maximizes scatter between classes and minimizes scatter within classes

Illustration of the Projection

- Using two classes as example:



Poor Projection




Good

LDA

We want a projection that maximizes:


$$J(w) = \max \frac{\textit{between class scatter}}{\textit{within class scatter}}$$

Fischer's Linear Discriminant Analysis



Slide inspired by N. Vasconcelos

Visualization

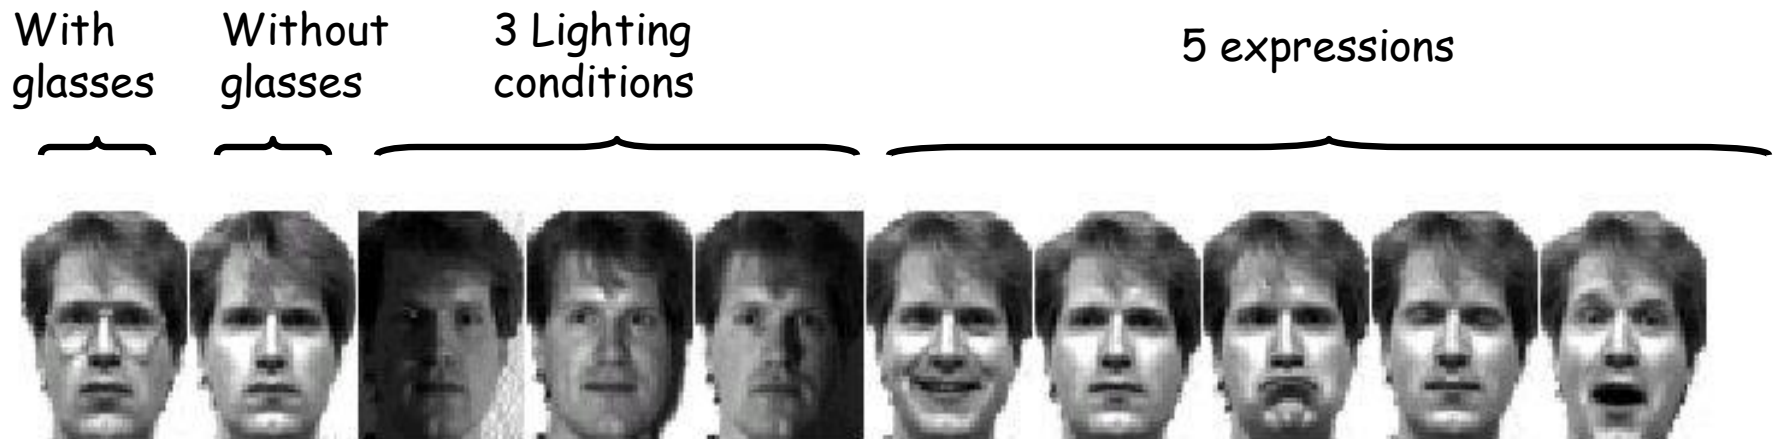


PCA vs. LDA


- Eigenfaces exploit the max scatter of the training images in face space
- Fisherfaces attempt to maximise the **between class scatter**, while minimising the **within class scatter**.

Results: Eigenface vs. Fisherface

- Input: 160 images of 16 people
- Train: 159 images
- Test: 1 image
- Variation in Facial Expression, Eyewear, and Lighting



Eigenface vs. Fisherface



What we have learned today

- Introduction to face recognition
- The Eigenfaces Algorithm
- Linear Discriminant Analysis (LDA)

Turk and Pentland, Eigenfaces for Recognition, *Journal of Cognitive Neuroscience* **3** (1): 71–86.

P. Belhumeur, J. Hespanha, and D. Kriegman. "Eigenfaces vs. Fisherfaces: Recognition Using Class Specific Linear Projection". *IEEE Transactions on pattern analysis and machine intelligence* **19** (7): 711. 1997.